

Software-as-a-Graph: Heterogeneous Graph Learning for Pre-Deployment Reliability and Dependability Analysis of Complex Distributed Systems

Abstract

Modern distributed software—publish–subscribe middleware and microservices—decouples components to achieve scalability, but that decoupling obscures how cascading failures propagate across message brokers, topics, shared libraries, and host nodes. Detecting systemically critical components *before* deployment is hard precisely because operational telemetry does not yet exist, and static code analysis cannot observe distributed communication topology. Uncontained cascades that reach production trigger costly downtime, emergency restart loops, and wasted energy.

We present **Software-as-a-Graph (SaG)**, an AI-driven framework for pre-deployment reliability analysis that represents a distributed architecture as a typed multigraph over Applications, Brokers, Topics, Execution Nodes, and Shared Libraries. Its primary instrument is a **predictive pathway**: a relation-specific **Heterogeneous Graph Transformer (HGT)** that learns multi-layer propagation patterns to forecast cascading failure blast radii, rank critical components, and score per-relationship criticality from Architecture-as-Code manifests alone. Because a rank says *where* to act but not *what to do*, SaG pairs the predictor with an **explanation layer**: an interpretable Reliability–Maintainability attribution model grounded in ISO/IEC 25010 and 25019 that diagnoses the structural root cause of a flagged component’s fragility. The two share no parameters, and only the predictor is scored against the simulation oracle.

We evaluate SaG on 1,770 components from seven synthetic scenarios and three real-world open-source systems (Autoware.universe ROS 2, GCP Cloud Microservices, Train-Ticket) against independent discrete-event cascade simulators under a strict input–label independence guarantee:

- **Typing decides transfer.** On unseen architectures the typed model reaches $\rho = 0.608$; relation-specific typing alone accounts for +0.353 over its homogeneous counterpart ($\rho = 0.439$ vs. 0.086), winning all eight cross-scenario folds ($p = 0.008$). QoS edge encoding adds +0.169 out of distribution (7/8 folds, $p = 0.016$).
- **Critical-set identification and CI/CD-viable cost.** The predictor identifies the critical top- K set better than any alternative evaluated ($F_1@K = 0.414$), and inference costs 44 ms on a 2,000-component system — two orders of magnitude below the structural analysis feeding it, making per-commit architectural gating practical.
- **An explicit boundary.** Against a training-free QoS-weighted structural baseline the ranking advantage is +0.037 and *not* significant ($p = 0.64$). We report this directly: graph learning earns its cost through transfer to unfamiliar architectures and through the typed attention, multi-task outputs, and relationship-level criticality it supplies alongside the ranking, not through ranking accuracy alone.
- **Real-world generalization.** SaG transfers zero-shot to authentic cyber-physical and cloud-native systems ($\rho = 0.688$ to 0.778), capturing up to 100% of failure-impactful services within the predicted critical set.

By enabling automated architectural analysis inside CI/CD before deployment, SaG bridges the Architecture–Code Gap, fostering resilient, dependable, and energy-efficient software systems.

Keywords: Graph representation learning, Heterogeneous graph neural networks, Publish–subscribe middleware, Distributed systems dependability, Cascading failures, Static system analysis, Architectural quality models, Explainable AI, CI/CD quality gates

1. Introduction

1.1. Motivation

Modern large-scale distributed systems increasingly rely on asynchronous, event-driven, and publish-subscribe (pub-sub) architectures. From autonomous vehicles (ROS 2 [1]) and high-throughput enterprise backbones (Apache Kafka [2]) to cyber-physical systems (DDS [3]) and IoT meshes (MQTT [4]), pub-sub decouples producers and consumers in space, time, and synchronization [5]. Components communicate indirectly by sending and receiving messages through intermediate topics and brokers without maintaining direct references to one another. Furthermore, modern middleware lets developers specify deployment-time Quality-of-Service (QoS) policies—such as message durability, transport reliability, and delivery deadlines—to control how data flows under stress.

While this decoupling makes distributed systems scalable and flexible, it introduces a severe **visibility barrier** for system reliability and dependability:

- **Indirect Failure Pathways:** In traditional synchronous systems (such as REST or RPC architectures), component interactions follow explicit caller-callee call graphs. In pub-sub and event-driven meshes, there are no direct static references between publishers and subscribers. Cascading failures propagate across hidden logical pathways spanning brokers, shared topics, colocated execution nodes, and shared software libraries.
- **Distinct Failure Mechanisms:** Failures in distributed systems do not propagate in a single way. They manifest as either *sequential cascades* (e.g., slow consumer message-queue backlogs propagating downstream) or *simultaneous blast radii* (e.g., a shared library crash or node failure instantly disabling multiple colocated services at the same moment). Traditional architectural diagrams and static call graphs fail to capture these multi-layer interactions.

Fixing these vulnerabilities is easiest and cheapest **prior to deployment**, during architectural design and Continuous Integration / Continuous Delivery (CI/CD). However, pre-deployment is precisely when **no runtime telemetry, distributed tracing, or operational logs exist**. As a result, software architects and reliability engineers face two fundamental questions without runtime data:

1. *Which components and communication links are systemically critical to system reliability and availability?*
2. *Why are they critical, and what specific architectural fix (such as replicating a broker, decoupling a shared topic, or sandboxing a library) will most effectively eliminate that risk?*

Addressing this challenge is also critical for **system performance and sustainability**. When cascading failures occur in production, they trigger compute-intensive restart loops, failover storms, and retransmissions. Proactive, design-time architectural hardening prevents these failures before software is deployed, saving infrastructure costs and eliminating wasted energy.

1.2. Problem Statement: The Architecture-Code Gap

We formulate pre-deployment dependability analysis around one primary predictive task and the explanation layer that makes its output actionable:

1. **Failure-Impact Forecasting (Predictive Pathway) — the primary task.** Forecasting the dynamic, global cascading failure blast radius and ranking the systemically critical component set, by training a data-driven, non-linear model over learned topological representations. No static aggregate of centrality measures supplies this: cascade reach is multi-hop and relation-dependent, and a component’s blast radius depends on *which* kind of edge carries the failure outward. It is also the only task here with an independent ground truth, and the only one we evaluate as a ranking model.
2. **Explainable Criticality Attribution (Explanation Layer) — what a rank alone cannot say.** A ranked set tells an architect *where* to act but not *what to do*. We therefore pair the predictor with an interpretable structural quality profile grounded in ISO/IEC 25010 [6] and ISO/IEC 25019 [7], which diagnoses the *qualitative nature* of a flagged component’s vulnerability—distinguishing, say, a single point of failure from a high-coupling maintenance bottleneck—and therefore which remediation applies. It is not a ranking model, and we do not evaluate it as one.

The separation is architectural, not merely presentational: the two pathways consume the same graph but share no parameters, and neither is fitted to the other’s output — the coupling term that would tie them is disabled by default and reported only as an ablation (Section 4.2). Keeping them distinct is what allows a component to be reported as, say, structurally central but operationally low-impact — a diagnosis neither pathway produces alone.

Existing software engineering approaches fail to bridge what we define as the “**Architecture–Code Gap**”: *a distributed system can have pristine, 100% bug-free source code within each service, yet remain fragile to catastrophic global outages due to hidden architectural single points of failure (SPOFs) or mismatched middleware QoS contracts*. Three prevailing paradigms leave this gap unaddressed:

- **Static Code Analysis (SCA)**: Tools like SonarQube inspect source code complexity [8], modularity, and cohesion [9, 10] within single services. However, SCA cannot see the broader distributed network, message queues, or cross-host failure propagation.
- **Runtime Chaos Engineering**: Techniques like Chaos Monkey [11] and distributed tracing inject real faults into running staging or production environments. While effective at validating live systems, they require fully deployed infrastructure, carry operational risks, and arrive too late to guide initial architectural design.
- **Homogeneous Graph Centrality Metrics**: Standard network metrics (betweenness, PageRank, degree) [12, 13, 14, 15] flatten systems into simple, unweighted graphs. They treat all connections identically, failing to distinguish between a message topic, a shared library, and an execution host. Similarly, homogeneous Graph Neural Networks (GNNs) [16, 17, 18] ignore relation-specific message routing.

Currently, no unified framework combines typed multigraph modeling, source-code quality metrics, heterogeneous graph learning, and explainable quality attribution for pre-deployment analysis.

1.3. The Software-as-a-Graph (SaG) Approach

To bridge this gap, we introduce **Software-as-a-Graph (SaG)**, an AI-driven pre-deployment **Static System Analysis (SSA)** framework. SaG ingests Architecture-as-Code manifests and executes a four-stage pipeline:

1. **Typed Multigraph Formulation**: SaG models the distributed architecture as a typed, directed multigraph over five core entity types: Applications, Brokers, Topics, Execution Nodes, and Shared Libraries (Section 3.1).
2. **QoS-Aware Logical Dependency Projection**: Using six formal projection rules, SaG derives a semantic `DEPENDS_ON` dependency layer that captures both sequential cascades (via topics and brokers) and simultaneous blast radii (via shared libraries and node colocation), weighted by declared QoS contracts (Section 3.2).
3. **Heterogeneous Graph Learning for Failure Forecasting (Predictive Pathway)**: SaG trains a **Heterogeneous Graph Transformer (HGT)** whose relation-specific attention lets a `USES` edge into a shared library propagate differently from a `PUBLISHES_TO` edge into a topic. It forecasts blast radii, ranks critical components, and emits per-relationship criticality and multi-task reliability/maintainability outputs (Section 4), at 44 ms per 2,000-component system.
4. **Explainable Quality Attribution (Explanation Layer)**: To say *why* a flagged component is critical, SaG combines code-level SCA metrics with topological properties into a deterministic **Reliability–Maintainability (RM)** attribution model (Section 5). Reliability decomposes into **Fault Tolerance** (error propagation depth) and **Availability** (single-point-of-failure exposure), which point at different repairs. Being a linear aggregate that models no propagation dynamics by construction, it explains *why* a component is vulnerable rather than how far a cascade travels; its standalone rank correlation is correspondingly modest ($\rho = 0.195$, Section 7.1) — a consequence of that scope, not a defect to tune away.

To ensure methodological rigor, SaG enforces an **input–label independence guarantee**: the learned models and attribution baseline operate strictly on the derived analytical graph G_{analysis} , whereas ground-truth failure impacts are generated by independent discrete-event simulators executing over the raw structural topology $G_{\text{structural}}$ (Section 4.4).




Figure 1: End-to-end architecture of the Software-as-a-Graph (SaG) framework. A shared front end (manifest ingestion → typed multigraph → QoS-weighted DEPENDS_ON projection → typed node features) feeds two deliberately separate pathways. **Pathway B (predictive)** is the primary one: it produces a ranked critical set and per-relationship criticality, and is the only pathway validated against the simulation oracle — the oracle scores rankings, which a quality profile is not, and it is strictly an offline training-and-validation component, never a dependency of online inference. **Pathway A (explanation)** then produces a standards-grounded quality profile for what B flagged, and the remediation it implies; it explains *why* a component is fragile and is not a ranking model. The single link between them is triage rather than data flow: the architect applies A’s explanation to whatever B flagged. The pathways share no parameters, and the oracle runs on $G_{\text{structural}}$ alone, never on the graph the predictors see (Section 4.4). A’s remediation guidance closes a loop of its own: each candidate edit is re-simulated on its own mutated copy of $G_{\text{structural}}$ and kept only if it beats the simulator’s own seed-to-seed noise, before being accepted.

Rationale for Graph Learning vs. Direct Simulation. Because discrete-event simulation $I^*(v)$ defines ground-truth criticality in this study, one might ask: *why train a graph neural network rather than running simulation sweeps directly?* If a complete simulator is available and compute budget is unlimited, simulation alone can identify critical components in an existing system. However, four practical reasons make our hybrid graph learning and attribution framework essential:

1. **Handling Unmeasured Components:** Discrete-event simulators only inject faults into active

application processes, leaving passive infrastructure (such as message topics or host nodes) without direct simulation labels (30% to 47% of components per system). The learned GNN generalizes across both labeled and unmeasured entities.

2. **Variance Reduction and CI/CD-Viable Cost:** Cascade simulations are noisy and highly sensitive to seeds and propagation thresholds (label standard deviation across seeds reaches 0.416); the network learns a smooth, threshold-marginalized representation instead. It is also two orders of magnitude cheaper than the analysis feeding it — 44 ms on a 2,000-component system against minutes to hours for exhaustive multi-seed simulation — which is what makes per-commit gating affordable rather than nightly (Section 7.5).
3. **Diagnostic Explainability:** Simulators and trained GNNs both return impact — a score, a rank, and in the GNN’s case an attention distribution showing *where* the model looked (Section 7.3). Neither returns cause attributed in standardised quality terms. A simulator can name precisely which subscriber lost which feed, and is fully inspectable in that sense, but it cannot say whether the component is fragile because it is a single point of failure or because it is a high-coupling maintenance bottleneck — and those two diagnoses call for different remediations (host or broker replication versus topic decoupling and refactoring). Our ISO-grounded RM model supplies exactly that missing layer: once the predictive path has identified a critical component, the diagnostic path says which quality characteristic is at risk and therefore which architectural fix applies.

1.4. Research Questions

Our empirical evaluation investigates five research questions:

RQ1 (Predictive Efficacy): *How accurately does heterogeneous graph learning predict cascading failure impact and identify the critical component set, compared to traditional, non-learning network metrics?*

RQ2 (Value of Architectural Typing): *Does modeling distinct entity and dependency types (applications, topics, brokers, hosts, and libraries) yield better failure predictions than homogeneous graph models, and does that advantage hold on architectures the model has never seen?*

RQ3 (QoS Encoding, Calibration and Sensitivity): *How do middleware Quality-of-Service policies, the declared weighting constants of the explanation layer, the choice of simulation oracle, and propagation-threshold and normalisation settings affect predictive accuracy, stability, and explainability?*

RQ4 (Real-World Generalization): *How effectively does the framework transfer zero-shot to real-world, open-source distributed systems across autonomous driving (ROS 2) and cloud-native microservice architectures?*

RQ5 (Analysis Cost and Deployment Feasibility): *What does pre-deployment analysis cost at CI/CD time, which pipeline stage dominates that cost, and does the resulting budget support per-commit architectural gating?*

1.5. Key Contributions

This paper makes four principal contributions:

1. **Heterogeneous Graph Learning for Pre-Deployment Dependability:** A relation-specific Heterogeneous Graph Transformer that forecasts cascading blast radii from Architecture-as-Code alone, with a 16-dimensional edge feature vector carrying 7 QoS dimensions and multi-task heads for component and relationship criticality (Section 4). Our central claim is about transfer: typing is what lets graph learning generalize to architectures it never trained on, where the untyped alternative collapses (Section 7.2).
2. **A Formal Typed Architecture Model:** A multigraph representation that derives logical dependencies from physical pub-sub linkage and distinguishes sequential cascade propagation from simultaneous multi-consumer library failures, supplying the typed substrate the predictor consumes (Section 3).

3. **A Standards-Grounded Explanation Layer:** An interpretable Reliability–Maintainability model grounded in ISO/IEC 25010/25019 that turns the predictor’s ranked output into an actionable diagnosis, separating single-point-of-failure exposure from error-propagation depth — two conditions calling for different repairs (Section 5).
4. **Empirical Evaluation, Real-World Validation, and Cost Characterisation:** A rigorous evaluation across seven synthetic topologies (1,545 components) and three real-world systems (Auto-ware.universe ROS 2, GCP Cloud Microservices, Train-Ticket; 225 components) under strict input–label independence, establishing both where typed graph learning delivers decisive advantages and the boundary where it does not, with a per-stage cost profile showing the learned model is the pipeline’s cheapest stage (Sections 6–7).

Relationship to the authors’ prior work. An earlier, shorter version was presented at a peer-reviewed conference [19], covering the typed multigraph formulation and the deterministic quality-attribution model on the synthetic corpus alone. This is a substantially extended version. New here: the entire predictive pathway as reported — the Heterogeneous Graph Transformer, its QoS edge encoding, the multi-task masked-loss heads, and every learned result in Section 7; the inductive LOSO protocol and cross-architecture transfer analysis carrying the paper’s central claim (Section 7.2); the three real-world validations (Section 7.4); the cost characterisation answering RQ5 (Section 7.5); the four-oracle convergent-validity analysis and independence guarantee (Sections 4.3–4.4); and the joint weight-sensitivity screening (Section 7.3). Carried-over material is confined to Sections 3 and 5 and is itself revised. No companion manuscript from this work is under review elsewhere.

1.6. Paper Organization

The remainder of this paper is organized as follows: Section 2 reviews related work in distributed systems dependability, static system analysis, and graph representation learning. Section 3 formalizes the Software-as-a-Graph architectural model, the dependency projection rules, and the typed node features both pathways consume. Section 4 presents the Heterogeneous Graph Transformer, its multi-task heads, the simulation oracles that supply its labels, and the input–label independence guarantee. Section 5 presents the interpretable RM explanation layer. Section 6 describes the experimental setup, benchmark corpus, and evaluation protocols. Section 7 presents empirical results for RQ1–RQ5. Section 8 discusses architectural implications, performance and sustainability implications, threats to validity, limitations, and concluding remarks.

2. Related Work

This research intersects four foundational domains: distributed systems dependability, static system analysis, graph neural networks for network vulnerability, and software quality models.

2.1. Dependability in Distributed and Pub-Sub Systems

The publish–subscribe (pub-sub) paradigm provides core communication decoupling for scalable distributed software [5]. Modern middleware standards—such as ROS 2 [1], Apache Kafka [2], DDS [3], and MQTT [4]—enable fine-grained Quality-of-Service (QoS) policies governing message durability, transport reliability, and queue deadlines. Prior dependability research in this area has focused primarily on **runtime mechanisms**, such as dynamic consensus protocols, broker clustering, adaptive retransmission, and automated failover.

In parallel, **chaos engineering and runtime verification** [11] inject simulated faults into running staging or production clusters to observe recovery behavior. While runtime fault injection is valuable for testing operational infrastructure, it comes with fundamental limitations:

- It requires fully deployed, operational testbeds.
- It carries operational risk if run near production.

- It operates too late in the software lifecycle to evaluate alternative architectural designs before systems are built.

Our work addresses the complementary, **pre-deployment phase**: predicting systemic cascading vulnerabilities directly from Architecture-as-Code descriptors *before* systems are deployed.

2.2. Static Code Analysis (SCA) vs. Static System Analysis (SSA)

Traditional **Static Code Analysis (SCA)** tools (such as SonarQube) inspect source code Abstract Syntax Trees (ASTs) within individual services. They evaluate cyclomatic complexity [8], class cohesion, module coupling (e.g., LCOM, CBO) [9, 10], and duplicated code to flag internal code smells and defect-prone components [20, 21, 22, 23]. However, SCA is completely blind to runtime topology: it cannot observe inter-service messaging channels, message broker queues, or cross-host container placement.

To close this “Architecture–Code Gap,” **Static System Analysis (SSA)** elevates static analysis from single-service source code to the global system architecture. By modeling distributed applications, message topics, brokers, and execution nodes as a connected graph, SSA propagates code-level metrics across architectural dependencies. This allows software teams to catch structural anti-patterns [24, 25, 26, 27] and architectural technical debt [28, 29] early during continuous integration (CI/CD) [30, 31].

2.3. Software Quality Models and Multi-Criteria Evaluation

Software product quality is standardized by the **ISO/IEC 25010:2023** product quality model [6] (which defines characteristics including Reliability, Maintainability, and Performance Efficiency) and the **ISO/IEC 25019:2023** Quality-in-Use model [7]. Software measurement distinguishes between *internal quality* (measured on static artifacts at rest) and *external quality* (measured on executing software) [32, 33].

Combining multi-attribute structural metrics into an overall quality score is a classic Multi-Criteria Decision Making (MCDM) problem. The **Analytic Hierarchy Process (AHP)** [34] provides a structured pairwise-comparison method with an explicit Consistency Ratio ($CR \leq 0.10$) to ensure weighting schemes remain mathematically sound. In this paper, we use AHP to construct an audited, explainable Reliability–Maintainability (RM) quality baseline, providing transparent architectural diagnostics alongside our learned graph models.

2.4. Network Science and Graph Representation Learning

Network science provides established centrality metrics to identify critical nodes, such as degree, closeness, betweenness centrality [12, 14], articulation points, and PageRank [13, 15]. Foundational studies on network robustness [35], cascading overloads [36], and interdependent networks [37] model how failures propagate across connected systems. However, standard network metrics suffer from two major limitations when applied to software architectures:

1. **Dimensional Collapse:** A single centrality number cannot distinguish *why* a component is critical—for instance, whether it is a single point of failure (SPOF), an error-propagating hub, or an over-shared library.
2. **Semantic Collapse:** Standard metrics treat all nodes and edges identically. They conflate fundamentally different architectural entities, such as an asynchronous message topic, a shared C++ library, and a physical execution host.

To move beyond hand-engineered graph metrics, recent research has applied machine learning to network vulnerability (e.g., FINDER [38], DrBC [39], and PowerGraph [40]). However, most existing models use **homogeneous message passing** (GCN [16], GraphSAGE [17], GAT [18]), which averages signals across all connections indiscriminately. Because distributed software architectures are inherently **heterogeneous** (comprising distinct entity types and relationship rules), homogeneous models blur critical architectural boundaries.

Heterogeneous Graph Neural Networks (RGCN [41], HAN [42], HGT [43], MAGNN [44]) solve this by using relation-specific transformations. We build upon the **Heterogeneous Graph Transformer (HGT)** architecture [43] to maintain typed relational semantics when forecasting cascading failure blast radii.

A learned ranker is actionable only if a practitioner can see what drove it. Attention weights are the most direct such signal in transformer-style architectures, and are what we inspect in Section 7.3; perturbation- and mask-based explainers such as GNNExplainer [45] instead search for the minimal subgraph preserving a prediction. Both families explain a model *in its own terms* — which nodes and edges mattered — and neither says *why* a flagged component is architecturally fragile in the vocabulary of a quality standard. That gap is what the explanation layer of Section 5 addresses.

Table 1: Comparison of dependability analysis paradigms for distributed systems.

Paradigm	Lifecycle Stage	Topology-Aware	Multi-Typed	Explainable	Zero-Runtime Needed
Static Code Analysis (SCA) [8, 9]	Pre-Deployment	No (Single Service)	No	Yes (Code Smells)	Yes
Chaos Engineering [11]	Post-Deployment	Yes (Live Cluster)	Partial	Partial (Logs)	No (Requires Cluster)
Network Centralities [12, 13]	Pre-Deployment	Yes (Flat Graph)	No	No (Single Scalar)	Yes
Homogeneous GNNs [16, 18]	Pre-Deployment	Yes (Flat Graph)	No	No (Black-Box)	Yes
Software-as-a-Graph (SaG)	Pre-Deployment	Yes (Multigraph)	Yes (5 Types)	Yes (ISO/IEC RM)	Yes (Manifest-Based)

3. The Software-as-a-Graph (SaG) Architectural Model

This section formalizes the Software-as-a-Graph multigraph representation (Section 3.1), the QoS-aware weighting and logical dependency derivation rules (Section 3.2), and the dual graph views utilized throughout the framework (Section 3.3).

3.1. Formal Multigraph Definition

We model a complex distributed system as a typed, weighted, directed multigraph:

$$\mathcal{G} = (V, E, \tau_V, \tau_E, w_V, w_E) \quad (1)$$

where:

- V is the set of system entities, partitioned into five disjoint types:

$$V = V_{\text{app}} \cup V_{\text{broker}} \cup V_{\text{topic}} \cup V_{\text{node}} \cup V_{\text{lib}} \quad (2)$$

- E is the set of directed edges connecting entities.
- $\tau_V : V \rightarrow \mathcal{T}_V$ and $\tau_E : E \rightarrow \mathcal{T}_E$ are typing functions assigning node and edge categories.
- $w_V : V \rightarrow [0, 1]$ and $w_E : E \rightarrow [0, 1]$ are weighting functions representing entity criticality and coupling strength.

Application and Library entities additionally ingest static code metrics computed via SCA tools (`cm_*` attributes: lines of code, cyclomatic complexity, coupling between objects, LCOM), directly bridging code-level fragility with topological analysis.

3.2. QoS-Aware Weights and Logical Dependency Derivation

In distributed middleware, communication links exhibit varying degrees of coupling based on their Quality-of-Service (QoS) contracts. For example, a `RELIABLE` topic with `TRANSIENT_LOCAL` durability binds communicating services far more tightly than a `BEST_EFFORT` telemetry stream.

Each topic t carries an intrinsic criticality weight $w(t) \in [0, 1]$ combining its declared QoS semantics with two runtime-stress modulators, payload size and publication frequency:

$$w(t) = \beta \cdot \text{QoS}(t) + \alpha \cdot \text{SizeNorm}(t) + \psi \cdot \text{FreqNorm}(t), \quad (\beta, \alpha, \psi) = (0.75, 0.15, 0.10) \quad (3)$$

where the QoS term is itself an AHP-weighted aggregate of the declared contract,

$$\text{QoS}(t) = w_{\text{rel}} \cdot q_{\text{rel}} + w_{\text{dur}} \cdot q_{\text{dur}} + w_{\text{prio}} \cdot q_{\text{prio}}, \quad (w_{\text{rel}}, w_{\text{dur}}, w_{\text{prio}}) = (0.24, 0.62, 0.14), \quad (4)$$

Table 2: Entity and structural edge types in the SaG model.

Entity Type ($\mathcal{T}_V$)	Architectural Role	Concrete System Examples
Application (V_{app})	Autonomous process producing/consuming messages	ROS 2 node, Kafka microservice, MQTT client
Broker (V_{broker})	Message routing and queuing intermediary	RabbitMQ exchange, Mosquitto, EMQX broker
Topic (V_{topic})	Named logical communication channel	<code>/sensor/lidar</code> , <code>orders.payment.completed</code>
Node (V_{node})	Physical host or virtualized environment	Bare-metal server, Kubernetes worker, Cloud V
Library (V_{lib})	Shared software package or driver dependency	<code>librdkafka</code> , OpenCV, Protobuf runtime
Structural Edge ($\mathcal{T}_E$)	Direction	Semantic Meaning
PUBLISHES_TO	App/Library $\rightarrow$ Topic	Component publishes messages to topic
SUBSCRIBES_TO	App/Library $\rightarrow$ Topic	Component consumes messages from topic
ROUTES	Broker $\rightarrow$ Topic	Broker manages and routes topic traffic
RUNS_ON	App/Broker $\rightarrow$ Node	Process is hosted on physical/virtual host
CONNECTS_TO	Node $\rightarrow$ Node	Physical network link between hosts
USES	App $\rightarrow$ Library	Application links to shared library dependency

with $q_{\text{rel}}, q_{\text{dur}}, q_{\text{prio}} \in [0, 1]$ the normalized reliability, durability and transport-priority scores. Durability dominates because it governs whether data survives at all (state persistence across restarts and network partitions), whereas reliability and transport priority both govern in-flight delivery quality, with reliability weighing somewhat more since an unconditional delivery guarantee precedes the scheduling of it. The sub-weight vector is the geometric-mean priority vector of an independently-stated Saaty pairwise-comparison matrix, with a small but genuinely nonzero consistency ratio ($CR \approx 0.016$). An earlier version of this matrix was solved backward from a chosen target vector rather than stated independently, which is why its near-zero CR was an artifact of construction and not evidence of a real judgement; we report the honest, independently-elicited CR here instead. The modulators are logarithmically compressed, $\text{SizeNorm}(t) = \log_2(1 + \text{bytes})/20$ (a 1 MiB design envelope, the practical DDS sample ceiling before RTPS fragmentation dominates) and $\text{FreqNorm}(t) = \log_{10}(1 + \text{Hz})/3$, and $w(t)$ is clamped to $[0.01, 1]$ so that best-effort edges remain visible to graph traversals. Every structural communication edge incident on t (PUBLISHES_TO, SUBSCRIBES_TO, ROUTES) inherits $w_E(e) = w(t)$ together with the topic’s QoS vector.

The outer split (β, α, ψ) is a declared convex combination rather than an elicited one, and we report its sensitivity directly (Section 7.3.2) rather than defending the particular triple. An earlier KiB-based SizeNorm divisor of 50 implied an unstated ~ 1 EiB envelope and left the term realizing only $\sim 2\%$ of alpha’s declared 15% budget on the evaluation corpus’s real payload sizes (32 B–32 KiB); the byte-based 1 MiB envelope above corrects that so alpha’s contribution is no longer negligible by construction.

3.2.1. Logical Dependency Projection (*DEPENDS_ON*)

Structural edges capture explicit deployment connections but omit implicit runtime dependencies. For example, a subscriber depends upon a publisher, yet no direct edge connects them in pub-sub architectures. We derive a single unified semantic relation, *DEPENDS_ON*, directed from *dependent* to *dependency* (“if target fails, source is impacted”):

Table 3: The six *DEPENDS_ON* logical dependency projection rules.

Rule	Dependency Category	Structural Pattern (Dependent $\rightarrow$ Dependency)	Derived Weight (w)
1	<code>app_to_app</code>	Subscriber $\rightarrow$ Publisher (via shared Topic, incl. transitive <i>USES</i>)	$1 - \prod_{t \in T} (1 - w(t))$
2	<code>app_to_broker</code>	Publisher/Subscriber $\rightarrow$ Broker routing its topics	$1 - \prod_{t \in T} (1 - w(t))$
3	<code>node_to_node</code>	Host $\rightarrow$ Host (lifted from inter-host app dependencies)	Lifted max w
4	<code>node_to_broker</code>	Host $\rightarrow$ Broker (lifted from hosted app dependencies)	Lifted max w
5	<code>app_to_lib</code>	Application $\rightarrow$ Shared Library it <i>USES</i>	$H(w_V(\text{app}), w_V(\text{lib}))$
6	<code>broker_to_broker</code>	Broker $\leftrightarrow$ Broker (shared-host fate, symmetric)	$w_V(\text{node})$

Rules 1 and 2 aggregate the several topics T mediating one component pair by *probabilistic union* rather than by a maximum, so that additional parallel failure vectors raise coupling monotonically while preserving $w \in (0, 1]$. Rule 5 uses the harmonic mean $H(x, y) = 2xy/(x + y)$ of the consuming Application’s and the shared Library’s own vertex weights, which calibrates caller criticality against dependency criticality instead of letting either endpoint dominate. Rules 3 and 4 lift the maximum weight of the component-level dependencies crossing the host boundary.

3.2.2. Sequential Cascades vs. Simultaneous Blasts

A key insight of the SaG model is distinguishing between two fundamentally different failure modes:

- **Sequential Cascade (Rule 1):** When an application publisher fails, downstream subscribers experience data starvation. The failure propagates step-by-step through topics and message queues.
- **Simultaneous Blast (Rule 5):** When a shared software library or execution node crashes, all consuming applications and colocated brokers fail *instantaneously* in a single event.

Retaining entity types and relation-specific dependency rules allows SaG to model both mechanisms, whereas untyped graphs collapse them into identical edges.

On the symmetry of Rule 6. Rule 6 is the one projection rule that is symmetric, and deliberately so: it does not assert that one broker functionally depends on another, but that two brokers co-located on the same host *share that host’s failure domain*. This is the same simultaneous-blast semantics as Rule 5, which is why the derived weight is the shared *Node’s* weight rather than either broker’s, and why the relation holds in both directions. The mechanism is real in deployed middleware — co-located brokers contend for CPU, page cache, file descriptors and NIC bandwidth, and a host loss takes both at once, which is precisely why operational guidance for Kafka, RabbitMQ and EMQX is to spread brokers across fault domains. What Rule 6 does *not* model is intra-cluster broker coupling proper (partition replication, controller or quorum election, federation and shovel links), which is directional and does not require co-location; extending the schema to express it is left to future work. We also note that the rule is close to inert on our corpus: it fires in four of the eight cached scenarios and contributes twelve directed edges in total across 1,770 components, and because the simulation oracles read only $G_{\text{structural}}$ (Section 4.4), it cannot influence any ground-truth label.

3.3. Dual Graph Views and Architectural Layers

The SaG framework maintains two distinct representations of the system:

1. **Structural Graph ($G_{\text{structural}}$):** The raw graph containing physical and structural edges (PUBLISHES_TO, ROUTES, RUNS_ON, USES). This view is consumed exclusively by discrete-event simulators to execute unbiased failure injections (Section 4.3).
2. **Analysis Graph (G_{analysis}):** The projected graph containing derived DEPENDS_ON edges annotated with QoS weights and ingested SCA code metrics. All GNN feature representations, graph embeddings, and analytical metrics are computed on G_{analysis} .

G_{analysis} is further organized into four analytical layers (Application, Middleware, Infrastructure, and Global System), allowing architects to evaluate criticality at subsystem levels (e.g., following MIL-STD-498 hierarchical structures).

3.4. Typed Node Feature Representation

Both pathways read the same typed node features off G_{analysis} : the predictor (Section 4) projects them per entity type before message passing, and the explanation layer (Section 5) aggregates them into its quality profile. SaG extracts feature vectors tailored to the 5 entity types:

- **Application ($|V_{\text{app}}|$, 23 dims):** Indices 0–17 represent shared topological metrics (in/out degree, betweenness, closeness, reverse PageRank, clustering coefficient, articulation score, bridge load). Indices 18–22 capture source code metrics extracted via Static Code Analysis (SCA): Lines of Code (LOC), Cyclomatic Complexity, Martin’s Instability metric ($I_{\text{code}} = \frac{C_e}{C_a + C_e}$), Lack of Cohesion in Methods (LCOM), and composite Code Quality Penalty (CQP).




Figure 2: Running example: the raw structural graph (left) and the `DEPENDS_ON` projection derived from it (right). The projection makes implicit runtime dependencies explicit — a subscriber depends on the publishers of its topics even though no structural edge joins them — while the simulators continue to operate on the structural view alone.

- **Library** ($|V_{\text{lib}}|$, **25 dims**): Same shared topological (0–17) and code quality (18–22) metrics as Application, plus two library-specific extras (indices 23–24): the size of the transitive reverse-USES closure (normalized) and the count of distinct subscribers reachable from that closure’s published topics (normalized) — the two structural drivers of a library’s blast radius under both simulators’ cascade rules that the code-quality metrics alone do not capture.
- **Broker** ($|V_{\text{broker}}|$, **19 dims**): Indices 0–17 shared topological metrics; index 18 represents normalized queue buffer capacity.
- **Topic** ($|V_{\text{topic}}|$, **22 dims**): Indices 0–17 shared topological metrics; indices 18–21 capture publisher count, subscriber count, log message frequency $\log(1 + \text{freq})$, and ordinal QoS criticality.
- **Infrastructure Node** ($|V_{\text{node}}|$, **20 dims**): Indices 0–17 shared topological metrics; indices 18–19 capture normalized CPU core allocation and physical memory (RAM).

4. Graph Learning for Failure-Impact Prediction

Cascading failure impact is non-linear, multi-hop and relation-dependent: how far an outage travels depends not on how many neighbours a component has but on *which kind* of edge carries the failure outward, and on what lies two and three hops beyond. No closed-form aggregate of centrality measures expresses that, which is why the primary pathway of Section 1.2 is a learned one. This section presents the HGT architecture and its typed edge representation (Section 4.1), the multi-task heads and dimension-masked loss (Section 4.2), the simulation oracles supplying its labels (Section 4.3), and the input–label independence guarantee keeping the two apart (Section 4.4).

4.1. Heterogeneous Graph Transformer Architecture

Because distributed systems comprise heterogeneous entity types (Applications, Libraries, Brokers, Topics, Infrastructure Nodes) and diverse interaction semantics (`PUBLISHES_TO`, `SUBSCRIBES_TO`, `RUNS_ON`, `CONNECTS_TO`, `USES`, etc.), we implement a 3-layer **Heterogeneous Graph Transformer (HGT)** architecture [43] with hidden dimension $D = 64$ and 4 attention heads.

4.1.1. Edge Feature Encoding (16-Dimensional)

To capture continuous QoS constraints and channel semantics, SaG encodes each directed edge $e = (u, v)$ as a 16-dimensional continuous-categorical vector $e_{uv} \in \mathbb{R}^{16}$:

- **Index 0:** Scalar QoS weight $w_E(e) \in (0, 1]$, the same edge weight defined in Section 3.2 (inherited from $w(t)$ for structural pub/sub edges; a noisy-OR union, worst-case lift, or harmonic mean of endpoint weights for derived DEPENDS_ON edges, per rule).
- **Index 1:** Normalized path count through edge e in G_{analysis} .
- **Indices 2–8:** 7-bit one-hot encoding for edge relationship types (PUBLISHES_TO, SUBSCRIBES_TO, ROUTES, RUNS_ON, CONNECTS_TO, USES, DEPENDS_ON).
- **Indices 9–15:** 7 explicit QoS profile dimensions, non-zero only for PUBLISHES_TO/SUBSCRIBES_TO edges (all other edge types receive zeros): (9) Reliability policy score (0.0 = best-effort, 1.0 = reliable); (10) Durability policy score (0.0 = volatile, 0.5 = transient local, 0.6 = transient, 1.0 = persistent); (11) Normalized message priority (0.0/0.33/0.66/1.0 for low/medium/high/urgent); (12) Deadline constraint active flag (0/1); (13) Log deadline $\log_{10}(1 + \text{deadline_ns}/10^6)$; (14) Log max blocking time $\log_{10}(1 + \text{max_blocking_ms})$; and (15) a QoS heterogeneity flag (1 if the edge’s reliability/durability/priority triple deviates from the scenario’s modal QoS profile, 0 otherwise).

The **EdgeFeatureEncoder** projects e_{uv} into the hidden space: $e'_{uv} = W_{\text{edge}}e_{uv}$. Prior to relational attention, the edge projection is added directly to the destination node representation: $\tilde{h}_v = h_v + e'_{uv}$.

4.1.2. Type-Specific Projection and Heterogeneous Message Passing

For each source node u and target node v with meta-relation $\tau(e) = (\tau(u), \phi(e), \tau(v))$:

1. **Type-Specific Projection:** Node feature vectors x_v (dimensions 19–23 depending on $\tau(v)$) are mapped into D -dimensional hidden space:

$$h_v^{(0)} = \text{LayerNorm}(\text{GELU}(W_{\tau(v)}x_v)) \quad (5)$$

2. **Relational Mutual Attention:** Query, Key, and Value projections compute relation-specific attention across heads:

$$\text{Attn}(u, e, v) = \text{Softmax}_{\forall u \in \mathcal{N}(v)} \left(\frac{K(u)W_{\text{att}, \phi(e)}Q(\tilde{v})^\top}{\sqrt{D/H}} \right) \quad (6)$$

$$\text{Msg}(u, e, v) = V(u)W_{\text{msg}, \phi(e)} \quad (7)$$

3. **Bidirectional Message Passing:** To capture downstream backpressure and upstream cascading failures, message passing is executed over both forward and reverse relations (G_{analysis} and $G_{\text{analysis}}^\top$).
4. **Residual Aggregation and Layer Normalization:**

$$h_v^{(l)} = \text{LayerNorm} \left(h_v^{(l-1)} + \text{Dropout} \left(\sum_{u \in \mathcal{N}(v)} \text{Attn}(u, e, v) \cdot \text{Msg}(u, e, v) \right) \right) \quad (8)$$

4.2. Multi-Task Prediction Heads and Dimension Masking

From the final node embeddings $h_v^{(L)}$, SaG branches into specialized multi-task prediction heads:

- **Reliability Head:** $\hat{R}(v) = \sigma(\text{MLP}_R(h_v)) \in [0, 1]$
- **Maintainability Head:** $\hat{M}(v) = \sigma(\text{MLP}_M(h_v)) \in [0, 1]$
- **Composite Failure Impact Head:** $\hat{I}^*(v) = \sigma(\text{MLP}_C(h_v \parallel \hat{R}(v) \parallel \hat{M}(v))) \in [0, 1]$
- **Relationship Criticality Head:** $\hat{Q}(u, v) = \sigma(\text{TypedEdgeEncoder}_{\phi(e)}(h_u, h_v, e_{uv})) \in [0, 1]$

4.2.1. Dimension-Masked Loss Formulation

The joint optimization objective balances regression accuracy, multi-task dimension learning, ranking fidelity, pairwise ordering, and edge prediction:

$$\mathcal{L} = \mathcal{L}_{\text{composite}} + 0.5 \cdot \mathcal{L}_{\text{dimension}} + 0.3 \cdot \mathcal{L}_{\text{rank}} + 0.1 \cdot \mathcal{L}_{\text{pairwise}} + 0.3 \cdot \mathcal{L}_{\text{edge}} + \lambda_{\text{RM}} \cdot \mathcal{L}_{\text{consistency}} \quad (9)$$

where $I^*(v)$ is the simulated cascade impact defined by the primary oracle of Section 4.3, $\mathcal{L}_{\text{composite}} = \text{MSE}(\hat{I}^*(v), I^*(v))$, $\mathcal{L}_{\text{rank}}$ is ListMLE ranking loss [46], $\mathcal{L}_{\text{pairwise}}$ is margin-ranking loss, and $\mathcal{L}_{\text{consistency}} = \text{MSE}([\hat{R}(v), \hat{M}(v)]_{v \in \text{unlabeled}}, [R_{\text{RM}}(v), M_{\text{RM}}(v)]_{v \in \text{unlabeled}})$ regresses the predicted R/M heads toward the diagnostic pathway’s own $R(v)/M(v)$ (Section 5), restricted to nodes the oracle never labeled. Headline results use $\lambda_{\text{RM}} = 0$: the two pathways are trained independently, matching the separation claimed in Section 1.2. Section 7 additionally reports $\lambda_{\text{RM}} = 0.1$ as an ablation, since a nonzero λ_{RM} is a real coupling between the pathways and not merely a regularizer on already-independent quantities.

Dimension Masking: Because dynamic cascade simulation ($I^*(v)$ via **FaultInjector**, the sole labeler used to train and evaluate every GNN result in this paper) observes runtime failure reachability rather than static code maintainability, maintainability ground truth is unobserved *under this labeler*. A separate oracle, $I_M(v)$ (Section 4.3), does compute a maintainability ground truth, but only at the Validate stage and only as an internal consistency check against the same dependency topology $M(v)$ is scored from — it is never substituted in as a training label, since doing so would not add the independent supervision the mask is meant to withhold. We introduce a boolean dimension mask $m = [m_R, m_M] = [1, 0]$:

$$\mathcal{L}_{\text{dimension}} = \frac{1}{\sum_d m_d} \sum_{d \in \{R, M\}} m_d \cdot \text{MSE}(\hat{d}(v), d^*(v)) \quad (10)$$

This prevents the unmeasured maintainability head from being penalized or driven toward zero by backpropagation.

4.2.2. Domain-Rewighted Criticality (Q_{domain})

To ground predictions in ISO/IEC 25019 Context of Use ($\vec{\omega} = [q_R, q_M]^\top$), the composite score is evaluated as:

$$Q_{\text{domain}}(v) = q_R \cdot \hat{R}(v) + q_M \cdot M_{\text{static}}(v) \quad (11)$$

where $M_{\text{static}}(v)$ is drawn directly from the structural analyzer’s maintainability baseline (y_{rm} maintainability column), combining learned dynamic reliability with static source-code maintainability. The implementation deliberately refuses to emit $Q_{\text{domain}}(v)$ unless the checkpoint’s maintainability head actually received real supervision — under the sole labeler used in every experiment reported here (**FaultInjector**, $m = [1, 0]$), maintainability is unmeasured, so $Q_{\text{domain}}(v)$ is never populated in these results. Domain reweighting sensitivity is instead evaluated directly against the static RM baseline, reported as a dedicated ablation (Section 7.3).

4.3. Ground-Truth Simulation Oracles

To evaluate predictive accuracy prior to deployment without relying on production runtime telemetry, SaG executes discrete-event failure simulations over the raw structural multigraph $G_{\text{structural}}$, complemented by one deterministic structural traversal for the dimension the simulations cannot observe. We establish a formal taxonomy of four component-level oracles and one relationship-level oracle:

- **Cascade Reachability Oracle ($I^*(v)$):** Implemented via **FaultInjector**, this oracle injects an unexpected node crash at component $v \in V$, propagates cascading failures across dependent topics, brokers, and network links using breadth-first dynamic traversal, and computes the fraction of surviving subscriber feeds severed by the outage. Publisher loss is weighted by publication rate, and the resulting per-topic feed loss is scaled by a QoS ladder ($\times 1.2$ for **RELIABLE**, $\times 1.15$ for high or urgent transport priority, $\times 1.05$ for medium) before being clamped to $[0, 1]$. The ladder does not read durability, which

carries the largest of the three QoS sub-weights (0.62; Section 3.2); we quantify how little this scaling moves the labels in Section 7.3. $I^*(v) \in [0, 1]$ serves as the primary continuous target label for training and evaluating GNN predictors.

- **Multi-Metric Composite Oracle ($I_{\text{comp}}(v)$):** Implemented via `FailureSimulator`, this oracle evaluates a multi-faceted failure impact vector:

$$I_{\text{comp}}(v) = 0.35 \cdot \Delta\text{Reachability} + 0.25 \cdot \Delta\text{Fragmentation} + 0.25 \cdot \Delta\text{Throughput} + 0.15 \cdot \Delta\text{FlowDisruption} \quad (12)$$

This is the most QoS-aware of the three: every term is weighted by an operational severity $s(t) = w(t) \cdot \text{rate}(t)$, so all three QoS dimensions enter through $w(t)$. Sixty per cent of the composite (0.35 reachability plus 0.25 fragmentation) nonetheless measures graph connectivity rather than message loss. $I_{\text{comp}}(v)$ serves as the canonical oracle for architectural quality gate verification.

- **Dynamic Queue-Flow Oracle ($I_{\text{dyn}}(v)$):** Implemented via `MessageFlowSimulator` using the SimPy discrete-event simulation framework [47], this oracle models message emission rates, stochastic network latencies, broker buffer saturation, and dropped delivery counts under fault injection. QoS policies shape per-message behaviour *inside* the run — reliability selects head-drop over tail-drop on queue overflow, and deadline and lifespan expiry discard samples — but the impact score extracted from it is the drop in delivered message rate suffered by surviving consumers, an unweighted per-copy count. $I_{\text{dyn}}(v)$ is therefore delivery-based and QoS-agnostic at the label level: a message on a safety-critical topic and one on a telemetry heartbeat contribute equally, and transport priority is not read at all.
- **Change-Propagation Oracle ($I_M(v)$):** Implemented via `ChangePropagationSimulator`, and the only oracle in this taxonomy that is not a discrete-event simulation: it is a deterministic breadth-first traversal of the *reversed* derived dependency graph (an edge $u \rightarrow v$ in G_{analysis} becomes $v \rightarrow u$ here, modelling “ v changing may force u to adapt”), stopped at loosely-coupled or interface-stable neighbours, and computed on every exhaustive Validate-stage sweep:

$$I_M(v) = 0.45 \cdot \text{ChangeReach}(v) + 0.35 \cdot \text{WeightedChangeImpact}(v) + 0.20 \cdot \text{NormalizedChangeDepth}(v) \quad (13)$$

$I_M(v)$ is the declared ground truth for the Maintainability dimension, but because it traverses the same derived dependency topology that $M(v)$ (Section 5.2) is itself scored from, $\rho(M, I_M)$ is an internal consistency check on that topology rather than independent behavioural evidence, unlike $I^*(v)$ or $I_{\text{dyn}}(v)$ against $R(v)$. It is never used as a GNN training label (Section 4.4).

- **Relationship (Edge) Removal Oracle ($I_{\text{edge}}(u, v)$):** Evaluates the systemic impact of severing an individual dependency or communication channel while keeping both endpoint components alive:

$$I_{\text{edge}}(u, v) = I_{\text{comp}}(G \setminus \{(u, v)\}) - I_{\text{comp}}(G) \quad (14)$$

Declared message criticality is withheld from the label, not from the features. `FailureSimulator` can additionally blend a declared `Topic.criticality` label into its severity term (a 50/50 blend with $w(t)$, via `use_topic_criticality`), and ships with this disabled. The reason is not merely that the label is author-supplied: `Topic.criticality` is *simultaneously* a GNN input feature (`topic_qos_criticality_ord`, live whenever the native substrate is QoS-embedded — the HGL-QoS arm and every LOSO fold), so an oracle that consumed it would score a predictor against a transform of its own input. In the synthetic corpus it is in any case not an independent operational judgement: it is `bucketize(w_QoS(t))` with 17% of labels flipped to a different level, a noised re-encoding of the same QoS score already inside $w(t)$ — the noise exists precisely so the label is not a deterministic function of the QoS features a model can see. Enabling the oracle flag would therefore only be sound as a dedicated ablation that simultaneously removes `topic_qos_criticality_ord` from the feature contract — which would invalidate every QoS-variant result in this paper and require a full retrain; we did not run that ablation.

Primary oracle declaration. Because the three reliability-facing component-level oracles (I^* , I_{comp} , I_{dyn}) measure different constructs and agree only partially (below), we designate exactly one of them as primary and hold every ranking claim in this paper to it. $I^*(v)$ (**FaultInjector**) is the primary oracle for all predictive-ranking results (Tables 7–9, RQ1–RQ3): it is the only oracle that both supplies training labels and scores held-out predictions, so evaluating a learned ranker against anything else would compare a model to a target it was never fitted to. The remaining oracles have narrower, explicitly named roles: $I_{\text{comp}}(v)$ (**FailureSimulator**) is the Validate-stage oracle for architectural quality gates and anti-pattern detection only, $I_{\text{dyn}}(v)$ (**MessageFlowSimulator**) is used only as an independent, behaviourally-constructed convergent-validity probe, $I_M(v)$ (**ChangePropagationSimulator**) is the declared but internally-consistent oracle for Maintainability alone and is not empirically validated in this paper (Section 8.3), and $I_{\text{edge}}(u, v)$ scores relationships rather than components. Each table and figure names the oracle it was measured against; no number is transferred between them.

These oracles are not interchangeable, and their disagreement bounds what any single one can support. Measured across seven scenarios and five seeds on the Application population, mean Spearman agreement is $\rho = 0.883$ for (I_{dyn}, I^*) , $\rho = 0.468$ for (I_{comp}, I^*) , and $\rho = 0.465$ for $(I_{\text{comp}}, I_{\text{dyn}})$, with per-scenario minima of 0.756, 0.171, and 0.121 respectively. The behavioural queue-flow oracle and the topological cascade oracle therefore agree strongly on *ordering*, which is genuine convergent evidence: they are constructed differently, one observing message delivery under load and the other reachability over edges, so their agreement is not reducible to a shared construction artifact. The composite oracle is the outlier, agreeing only moderately with either.

Agreement on the *critical set* is markedly weaker than agreement on ordering, and this is the binding limitation. Mean top- K Jaccard overlap at $K = 0.2n$ is 0.42 for the strongest pair and 0.31 for the other two, against 0.111 expected under independent rankings — around three to four times chance, but far from set identity. Two controls establish what this does and does not mean. First, it is not an artifact of tie-breaking: recomputing the overlap so that every component tied with the K -th is admitted changes it by at most 0.005. Second, it is not simply label noise, but neither is it purely construct divergence — the labeler’s own seed-to-seed self-agreement spans top- K Jaccard 0.44–1.00 (test–retest $\rho = 0.81$ –1.00), so a single oracle compared against itself reaches only 0.44 in its worst scenario. Top- K set identity is an intrinsically unstable statistic at this K , and the cross-oracle values sit just below the bottom of the range one oracle achieves against itself.

The consequence for how our results should be read is unchanged: a result established against one oracle is not evidence for a claim measured against another, and the instruction to “simply simulate” is under-specified — it does not say which simulator, at which propagation threshold, under which seed, and the answer materially changes the critical set. We report this disagreement as a bound on our own construct validity rather than as corroboration, and we note that we could not reduce it by weighting: applying the same $w(t)$ to both topological oracles does not make them converge (Section 7.3). That null result should be read as bounding the effect of the particular treatment we could apply, not of QoS-aware labelling in general: disabling QoS zeroes I_{comp} ’s severity term entirely, but on I^* it only removes a ladder multiplier that a $[0, 1]$ clamp already suppresses whenever a component is at or near total feed loss, so the two arms are not comparably strong (Section 7.3).

4.4. Input–Label Independence Guarantee

To guarantee experimental rigor and eliminate data leakage, SaG enforces strict architectural decoupling:

- **Feature Space:** Derived exclusively from G_{analysis} (using static structural topology, SCA metrics, and declared QoS contracts).
- **Label Space:** Evaluated exclusively on raw $G_{\text{structural}}$ through independent simulation oracles (**FaultInjector**, **FailureSimulator**, **MessageFlowSimulator**), plus **ChangePropagationSimulator** over the derived dependency view for the Maintainability dimension alone (Section 4.3).

No simulation outputs are ever exposed as input features to the GNN or attribution baseline.

5. The Explanation Layer: Standards-Grounded Criticality Attribution

The predictor of Section 4 answers *where* to act. It does not answer *what to do*, and neither does the oracle that scores it — both return impact, not cause attributed in standardised quality terms. A component may be critical because it is a single point of failure, because it propagates errors widely, or because it is a high-coupling maintenance bottleneck, and those diagnoses call for different repairs: replicate the host or broker, decouple the topic, or refactor the module. This section presents the layer supplying that step. SaG decomposes component and relationship criticality into a standards-grounded quality profile, computed over the same typed node features (Section 3.4) but sharing no parameters with the predictor, and applied to whatever the predictor flagged — triage rather than data flow (Figure 1).

5.1. Grounding in ISO/IEC Standards

Following **ISO/IEC 25010:2023** (Product Quality Model) [6] and **ISO/IEC 25019:2023** (Quality-in-Use) [7], we formulate two core formal criticality constructs:

- **Component Criticality (D_1):** The degree to which the sudden failure, unexpected termination, or severe degradation of an individual component reduces the system’s capacity to deliver required services within its operational context of use.
- **Relationship Criticality (D_2):** The degree of systemic service degradation resulting from the severance, partitioning, or failure of a specific dependency or communication channel while both endpoint components remain operational.

Criticality is evaluated primarily across two orthogonal quality characteristics: **Reliability (R)** and **Maintainability (M)**.

Table 4: The Reliability–Maintainability (RM) quality decomposition.

Dimension	Sub-Characteristic	Architectural Question	Underlying Graph Metrics
Reliability (R)	Fault Tolerance (FT)	How broadly does failure propagate?	Reverse PageRank on $G^\top$, in-deg
	Availability (A)	Is this a single point of failure?	Directed articulation score (raw)
Maintainability (M)	Modularity/Modifiability	How complex and coupled is this?	Betweenness, QoS-weighted out-

Coverage Scope: SaG focuses on Reliability and Maintainability. Safety (which requires domain-specific hazard logs, e.g., ISO 26262 ASIL ratings) and Security (which requires explicit threat models, e.g., STRIDE) are declared external to purely structural analysis and are left for domain-specific extensions.

5.2. Composite Quality Score Formulation

All raw topological and code metrics are rank-normalized to $[0, 1]$. The quality sub-characteristics are formulated hierarchically using the Analytic Hierarchy Process (AHP) [34]:

1. **Fault Tolerance ($FT(v)$):** Measures error cascade potential on the transpose graph $G_{\text{analysis}}^\top$ (where edges follow failure propagation from dependency to dependent):

$$FT(v) = 0.45 \cdot \text{RPR}(v) + 0.30 \cdot \text{Deg}_{\text{in}}(v) + 0.25 \cdot \text{CDPot}_{\text{enh}}(v) \quad (15)$$

where RPR is Reverse PageRank and $\text{CDPot}_{\text{enh}}$ is an enhanced Cascade Depth Potential term.

2. **Availability ($A(v)$):** Identifies structural single points of failure (SPOFs) across five terms — directed articulation severity, its QoS-weighted variant, edge-level irrecoverability, connectivity degradation, and the component’s own QoS weight:

$$A(v) = 0.35 \cdot \text{AP}_c^{\text{dir}}(v) + 0.25 \cdot \text{QSPOF}(v) + 0.25 \cdot \text{BR}(v) + 0.10 \cdot \text{CDI}(v) + 0.05 \cdot w(v) \quad (16)$$

3. **Reliability** ($R(v)$): Blends Fault Tolerance and Availability hierarchically:

$$R(v) = \alpha \cdot FT(v) + (1 - \alpha) \cdot A(v), \quad \alpha = 0.36 \quad (17)$$

Intra-dimension pairwise comparison matrices are audited against Saaty’s consistency ratio and measure $CR = 0.001$ (Fault Tolerance), $CR = 0.001$ (Availability), and $CR = 0.000$ (Maintainability) — all well within the $CR \leq 0.10$ acceptability threshold. The shipped intra-dimension weights are a $\lambda = 0.70$ shrinkage blend between the raw AHP-derived vector and a uniform prior (Section 7.3 reports the sensitivity of ranking accuracy to λ).

4. **Maintainability** ($M(v)$): Evaluates structural coupling combined with code-level static analysis across five terms — betweenness, QoS-weighted efferent coupling, the Code Quality Penalty, an afferent/efferent coupling-risk imbalance term, and inverse clustering:

$$M(v) = 0.35 \cdot BT(v) + 0.30 \cdot w_{\text{out}}(v) + 0.15 \cdot CQP(v) + 0.12 \cdot \text{CouplingRisk}_{\text{enh}}(v) + 0.08 \cdot (1 - CC(v)) \quad (18)$$

The baseline composite quality score $Q(v)$ combines both dimensions:

$$Q(v) = 0.80 \cdot R(v) + 0.20 \cdot M(v) \quad (19)$$

When evaluating under a specific ISO/IEC 25019 Context of Use vector $\vec{\omega} = [q_R, q_M]^\top$, the score is reweighted dynamically:

$$Q_{\text{domain}}(v) = q_R \cdot R(v) + q_M \cdot M_{\text{static}}(v) \quad (20)$$

Components are categorized into adaptive criticality tiers using box-plot quartile thresholds:

- **CRITICAL:** $Q > Q_3 + 1.5 \cdot \text{IQR}$
- **HIGH:** $Q_3 < Q \leq Q_3 + 1.5 \cdot \text{IQR}$
- **MEDIUM:** $Q_1 < Q \leq Q_3$
- **MINIMAL:** $Q \leq Q_1$

This enables actionable diagnostics: a service scoring high on A but low on FT is diagnosed as a pure SPOF requiring horizontal replication, whereas a service scoring high on FT is an error cascade hub requiring circuit breakers and bulkhead isolation.

6. Experimental Setup

6.1. Datasets and System Corpus

Our evaluation corpus comprises 1,770 components across ten distributed system architectures: $|V|$ is the sum of all five entity-type counts per scenario and sums to exactly the corpus total stated above. $|E|$ counts every raw structural relationship instance recorded in the scenario file (PUBLISHES_TO, SUBSCRIBES_TO, ROUTES, RUNS_ON, CONNECTS_TO, USES) — the substrate the simulation oracles traverse — not the denser derived DEPENDS_ON projection used for GNN training.

The three real-world architectures are hand-transcribed from open-source repositories using dedicated architectural adapters. The seven synthetic scenarios are produced by a parameterised topology generator: each is fully specified by a committed configuration giving a random seed, per-entity-type counts, seven-number summaries (mean, median, standard deviation, min, max, Q_1 , Q_3) for application publish and subscribe fan-out, for applications per host, for library fan-in and for topic payload size, and categorical distributions over the three QoS dimensions. Degree distribution and clustering are *emergent* from those inputs rather than specified directly. Table 6 gives the parameters that determine each topology’s shape; the complete configurations are in the replication package.

QoS diversity is not uniform across the corpus. Per-topic QoS is assigned by a domain-keyed lookup (`get_qos_for_topic`) that takes precedence over the categorical distribution declared in each scenario’s

Table 5: Experimental evaluation corpus.

Dataset / Architecture	System Paradigm	$ V $	$ V_{\text{app}} $	Topics	Brokers	Hosts	Libs	$ E $
Autonomous Vehicle (AV)	ROS 2 Cyber-Physical	152	80	40	4	8	20	797
Enterprise Pub-Sub	Kafka Event Mesh	520	300	120	10	40	50	3,245
Financial Trading	Low-Latency Pub-Sub	124	60	35	5	6	18	580
Healthcare Integration	HL7/FHIR Event Mesh	98	50	25	3	8	12	400
Hub-and-Spoke Enterprise	Broker-Centric Messaging	139	70	30	2	12	25	797
IoT Smart City	MQTT Telemetry Mesh	326	200	80	6	30	10	1,322
Microservices Mesh	Cloud-Native Services	186	90	45	6	15	30	680
Autoware.universe [48]	Real-World ROS 2 Autoware	75	32	24	3	6	10	179
Cloud Microservices [49]	Real-World GCP Boutique	60	22	20	4	6	8	128
Train-Ticket [50]	Real-World Microservices	90	41	30	3	8	8	162
Total		1,770						

Table 6: Generative parameters of the seven synthetic evaluation scenarios. Counts, seed and fan-out figures are read directly from the committed configurations. The modal QoS column gives the most common reliability/durability/priority value and the range of topic shares carrying them, computed from the committed topology rather than the config’s declared QoS targets, which domain-driven assignment does not always realize (Section 6.1).

Scenario	Config	Seed	Counts	Pub	Sub	Modal QoS (T)
Autonomous Vehicle (AV)	scenario_01_autonomous_vehicle	1001	80/40/4/8/20	2.5	5.0	RELIABLE/T
Enterprise Pub-Sub	scenario_07_enterprise_xlarge	7007	300/120/10/40/50	3.0	4.5	RELIABLE/T
Financial Trading	scenario_03_financial_trading	3003	60/35/5/6/18	4.0	6.0	RELIABLE/P
Healthcare Integration	scenario_04_healthcare	4004	50/25/3/8/12	2.5	3.0	RELIABLE/P
Hub-and-Spoke	scenario_05_hub_and_spoke	5005	70/30/2/12/25	2.0	7.0	RELIABLE/T
IoT Smart City	scenario_02_iot_smart_city	2002	200/80/6/30/10	2.0	1.5	BEST_EFFO
Microservices Mesh	scenario_06_microservices	6006	90/45/6/15/30	1.5	2.0	RELIABLE/T

configuration whenever a domain is set — true of every scenario here. For three domains (**hub-and-spoke**, **microservices**, **enterprise**) that lookup table has a single entry, so every topic receives an identical (reliability, durability, priority) triple regardless of name; a fourth (**av**) collapses four of five entries to the same triple. Table 6’s Modal QoS column reflects this directly: those four scenarios show 85–100% topic share at the single modal triple, against 51–79% for the three domains (**finance**, **healthcare**, **iot**) whose lookup tables are genuinely multi-valued. Concretely, $w(t)$ has standard deviation ≈ 0.02 on a $[0, 1]$ scale in the four QoS-flat scenarios, versus 0.19–0.28 in the other three. We report this as a corpus limitation rather than correct it: doing so would change what those four scenarios generate and require regenerating every cached result built on them (Section 8.3).

A note on the ATM case study. One further scenario, an Automated Teller Machine network, serves as an eighth LOSO fold (Section 6.3) and as the subject of the attention analysis (Section 7.3), but is deliberately *not* a row in Table 5 and contributes none of the 1,770 components counted there: it was authored as an illustrative walkthrough, and its configuration lacks the seven-number fan-out summaries the other synthetic scenarios declare. It enters LOSO because that protocol needs held-out topologies rather than characterised ones, so every LOSO figure reports eight folds against a ten-architecture corpus.

6.1.1. Reproducibility of the Corpus

The corpus is regenerable rather than merely archived. Each dataset is produced from its configuration by

```
python cli/generate_graph.py batch -input-dir data/scenarios -output-dir <dir>
```

and a manifest records, per dataset, the seed, the entity counts, the generating commit, and a SHA-256 digest of the emitted topology. A regression test asserts that every committed dataset regenerates *byte-identically* from its configuration and that the manifest matches what is on disk, so a divergence between the published numbers and the published data fails the test suite rather than passing unnoticed. A third party can therefore reproduce the exact graphs these results were computed on, not merely graphs drawn from the same distribution.

6.2. Baselines and Evaluated Predictors

We compare four primary predictor configurations, listing the learned predictors proposed in this paper first and the training-free baselines they are measured against second:

1. **HGL / HGL-QoS** (proposed): Relation-specific Heterogeneous Graph Transformers (Section 4), without and with the explicit 7-dimensional QoS edge encoding.
2. **GL / GL-QoS** (proposed ablation): Homogeneous Graph Attention Networks (GAT) [18] trained on the flattened, untyped graph projection — the controlled comparison that isolates what relation-specific typing buys.
3. **Topo-QoS** (baseline): QoS-weighted topological centrality baseline, training-free.
4. **Topo-BL** (baseline): Unweighted structural centrality (betweenness centrality and articulation point scoring), training-free.

The out-of-distribution table (Table 9) additionally reports **RM** / $Q(v)$ (the deterministic hierarchical quality attribution model of Section 5) as a non-competing diagnostic reference point, not as a fifth predictor: RM is not fitted to rank components, and its row exists to show how much the predictive path adds over static attribution (Section 1.2), not to compete on ranking accuracy. The same deterministic RM scorer is also the instrument behind every sensitivity sweep in Section 7.3: because RM is a closed-form function of its declared constants, sweeping those constants isolates their effect from training variance in a way the learned predictors cannot.

6.2.1. Evaluation Substrate

Predictors in this comparison are not all evaluated over the same graph view, and the distinction matters for reading the tables that follow. **Topo-BL, Topo-QoS, GL and GL-QoS** are scored on the derived Application–Library `DEPENDS_ON` projection (Section 3.2) — the same substrate built for GNN feature/label alignment and discussed further in Section 7.3 — while **HGL and HGL-QoS** consume the full native typed multigraph over all five entity types. **No predictor in either group reads $G_{\text{structural}}$** : that graph remains the exclusive substrate of the ground-truth simulation oracles (Section 4.4), and the independence guarantee enforced by `tests/test_independence_guarantee.py` holds identically for every variant in this section.

The projection substrate was adopted for the homogeneous and untrained baselines because the full native pub-sub graph is not learnable by an untyped model on this corpus: Application nodes carry near-zero betweenness and bridge-ratio in the raw graph (they never route messages), producing degenerate, near-constant feature vectors, and a single homogeneous message-passing layer over-smooths every Application into an identical representation by aggregating the same high-fan-in Topic hub. This is the design rationale recorded in the replication package rather than a measured ablation — we did not run GL/GL-QoS on the native substrate to quantify the collapse directly, and do not report a number for it.

This asymmetry is a scope condition on the RQ2 comparisons in Section 7.2: part of HGL’s advantage over GL reflects the wider native node set the typed model can consume, not relation-specific typing alone, since the untyped baseline could not be evaluated on that same view. It does not, however, affect which nodes are scored: the evaluation population is pinned identically across every variant, resolved from the native graph and the simulation labels alone — never from a variant’s own substrate or predictions — so no comparison in this paper is confounded by which nodes a particular predictor happened to see (Section 6.3).

6.3. Evaluation Metrics and Protocols

- **Ranking Accuracy:** Spearman rank correlation (ρ) and Kendall’s tau (τ) between predicted rankings and ground-truth simulated impact $I^*(v)$, the primary oracle declared in Section 4.3.
- **Critical-Set Identification:** $F_1@K$, Precision@ K , and Recall@ K for top- K critical components, where K is 20% of the evaluated node population, rounded to nearest. Because the predicted and true sets both contain exactly K elements, precision, recall and F_1 coincide at K ; the figure is a top- K overlap and we read it as such.
- **Statistical Significance:** Paired Wilcoxon signed-rank tests [51] ($p < 0.05$) and bootstrap 95% confidence intervals ($B = 2,000$) [52, 53].

6.3.1. Evaluation Population

Every predictor in a given table is scored on an identical node set, resolved from the graph and the labels alone — never from any variant’s predictions — so that no comparison is confounded by which nodes a particular method happened to score. Unless stated otherwise that set is the **Application** population: it is the population the framework’s central claim is about (topology predicts application-layer cascade criticality), and it is the only one every variant can score. This matters more than it may appear. Pooling node types into a single correlation mixes populations with different impact scales and base rates, and on this corpus that pooling is not benign — it moves the RM composite’s rank correlation *outside* the range spanned by its own per-type correlations (Section 7.3). We therefore report stratified, single-population figures throughout, and flag explicitly wherever a pooled figure appears.

6.3.2. Evaluation Protocols

- **In-Distribution Evaluation:** 60% train / 20% validation / 20% test node splits pinned by node identity within each scenario, evaluated over five random seeds {42, 123, 456, 789, 2024}.
- **Inductive Leave-One-Scenario-Out (LOSO):** Across all eight cached scenarios (the seven synthetic scenarios plus the ATM case study, Section 7.3), models are trained on the remaining seven and evaluated zero-shot on the held-out scenario, for eight folds total, to test out-of-distribution generalizability.

- **Real-World Architectural Transfer:** Evaluating zero-shot transfer on authentic open-source architectures.

7. Results and Empirical Analysis

7.1. RQ1: Graph Learning vs. Structural Baselines

Table 7 presents the in-distribution held-out Spearman rank correlation (ρ) against simulated cascade impact $I^*(v)$ across all seven synthetic scenarios.

Table 7: In-distribution held-out Spearman rank correlation (ρ) against $I^*(v)$ (seed means over 5 seeds with bootstrap 95% CIs; n is held-out Application count).

Scenario	n	Topo-BL	Topo-QoS	GL	GL-QoS	HGL	HGL-QoS
AV System	16	0.283	0.701	0.764	0.678	0.789	0.649
Enterprise	60	0.420	0.789	0.871	0.483	0.880	0.891
Financial Trading	12	0.289	0.700	0.645	0.539	0.854	0.770
Healthcare	10	-0.101	0.772	0.623	0.463	0.652	0.645
Hub-and-Spoke	14	0.234	0.359	0.547	0.054	0.534	0.297
IoT Smart City	40	-0.063	0.073	0.289	0.291	0.881	0.842
Microservices	18	0.401	0.707	0.525	0.564	0.483	0.476
Mean	—	0.209	0.586	0.609	0.439	0.725	0.653

Table 8: Paired Wilcoxon signed-rank tests across scenarios ($n = 7$, two-sided).

Comparison	$\Delta\rho$	Won	Wilcoxon W	p -value	Significance
HGL vs. Topo-BL	+0.516	7/7	0.0	0.0156	Statistically Significant ($p < 0.05$)
HGL vs. GL-QoS	+0.286	6/7	1.0	0.0312	Statistically Significant ($p < 0.05$)
HGL vs. Topo-QoS	+0.139	5/7	9.0	0.469	Not significant
GL vs. Topo-QoS	+0.023	4/7	10.0	0.578	Not significant
HGL vs. GL	+0.116	5/7	7.0	0.297	Not significant
HGL-QoS vs. HGL	-0.072	1/7	3.0	0.078	Not significant (marginal; HGL-QoS trails in-distribution)

7.1.1. Out-of-Distribution (LOSO) Generalization

Under inductive Leave-One-Scenario-Out (LOSO) cross-validation, the model must predict cascading criticality on an unseen system topology:

Eight LOSO folds are reported (the seven synthetic scenarios plus the ATM case study of Section 7.3), each holding out one scenario and training on the remaining seven. Every variant is scored on the identical Application node set per fold (Section 6.3); paired Wilcoxon tests below are over the eight folds.

Key Insights for RQ1:

1. **The typed predictor is the best configuration overall, and its margin over every learned alternative is decisive.** HGL-QoS leads on both metrics ($\rho = 0.608$, $F_1@K = 0.414$), with a large, significant margin over both homogeneous GNNs (+0.246 over GL-QoS, 8/8 folds, $p = 0.0078$).
2. **Critical-set detection is where the learned model separates from the untrained baselines.** HGL-QoS improves $F_1@K$ over Topo-QoS by +0.034 (0.414 vs. 0.380, a 8.9% relative gain) and over GL by +0.177. The advantage is real but far smaller than a pooled-population reading of the same experiment suggests.
3. **QoS encoding is what carries the typed model out of distribution.** HGL-QoS leads HGL by $\Delta\rho = +0.169$, winning 7 of 8 folds ($p = 0.0156$) — a larger and more consistent effect than the in-distribution comparison in Table 8 suggests in the opposite direction. See Section 7.3.

Table 9: Inductive Leave-One-Scenario-Out (LOSO) evaluation, Application population, eight folds. Rows are grouped by role, not listed as competing variants: training-free structural baselines, the learned predictors proposed in this paper, and the RM/ $Q(v)$ diagnostic reference of Section 1.2 (Pathway A), which is not a ranking model (Section 5.1) and is included only to quantify how much the predictive path adds over static attribution.

Predictor / Reference	Mean LOSO ρ	Std ρ	Critical-Set $F_1@K$	Requires Training
<i>Training-free structural baselines</i>				
Topo-BL	0.301	0.126	0.363	No
Topo-QoS	0.571	0.181	0.380	No
<i>Learned predictors (Pathway B)</i>				
GL (Homogeneous)	0.086	0.122	0.237	Yes
GL-QoS (Homogeneous)	0.363	0.089	0.341	Yes
HGL (Typed Heterogeneous)	0.439	0.145	0.327	Yes
HGL-QoS (Typed + QoS)	0.608	0.143	0.414	Yes
<i>Diagnostic reference (Pathway A) — not a ranking model</i>				
RM / $Q(v)$	0.195	0.130	0.327	No

- The boundary: on ranking alone, the untrained QoS baseline is not beaten.** Against *Topo-QoS*, HGL-QoS's margin is +0.037, won in only 5 of 8 folds, and **not statistically significant** ($p = 0.64$). We state this plainly: on out-of-distribution Application ranking we cannot claim heterogeneous graph learning outperforms a well-constructed structural baseline, only that it matches it. The case for the predictive pathway rests on the three findings above and on what it supplies that a centrality score cannot — typed attention, per-relationship criticality, multi-task quality outputs — not on ranking accuracy alone.
- The explanation layer is weakly predictive, not anti-predictive.** RM/ $Q(v)$ reaches $\rho = 0.195$ under distribution shift — better than untyped GL, worse than every QoS-aware variant. Its value is interpretable attribution rather than ranking (Section 5), but it is not noise.

A note on population. An earlier version of this table pooled every node type carrying a simulated label, which inverted several conclusions: it placed the RM baseline at $\rho = -0.014$ rather than +0.195 and GL at 0.381 rather than 0.086. This is the Simpson's-paradox hazard analysed in Section 8.3, and it is why every figure here is reported on a single, stated population.




Figure 3: Results at a glance, all on the Application population. (A) Out-of-distribution rank correlation per variant, whiskers showing σ across the eight LOSO folds; note that the training-free Topo-QoS baseline places second. The hatched bar is RM/ $Q(v)$, the Pathway A diagnostic reference of Section 1.2 — shown for context, not as a competing ranking predictor. (B) Critical-set detection at $K = 20\%$, same hatching convention. (C) Agreement between the three simulation oracles, whiskers showing the observed range across scenarios; the annotation gives top- K set agreement against both chance and the labeler's own seed-to-seed floor.

7.2. RQ2: Value of Typed Heterogeneity

To evaluate the contribution of node and edge typing, we compare relation-specific HGL against homogeneous GL across different validation regimes:

- **In-Distribution:** Heterogeneous HGL leads homogeneous GL by $\Delta\rho = +0.116$ (0.725 vs. 0.609; not statistically significant, $p = 0.297$, 5/7 scenarios won). On familiar topologies, homogeneous GNNs can partially approximate type distinctions through structural degree signatures, so the typed model holds only a point-estimate edge. The contrast with the out-of-distribution result below is the finding: typing buys little when the topology is already known and a great deal when it is not.
- **Out-of-Distribution (LOSO):** This is where typing decides the outcome. Heterogeneous HGL outperforms homogeneous GL by $+0.353$ ($\rho = 0.439$ vs. 0.086), winning *all eight* folds (paired Wilcoxon, $p = 0.0078$), and the gap between the QoS-aware variants is $+0.246$ (HGL-QoS 0.608 vs. GL-QoS 0.363), also 8/8 and $p = 0.0078$. The untyped model very nearly fails to transfer at all on Application ranking: at $\rho = 0.086$ it is below even the unweighted structural baseline (0.301). When encountering unseen topologies, relation-specific message passing is not an incremental refinement but the difference between a model that generalizes and one that does not.

A scope condition on this comparison. GL and HGL are not evaluated on the same graph view: as declared in Section 6.2.1, GL/GL-QoS run on the Application–Library `DEPENDS_ON` projection because the homogeneous model collapses to constant output on the full native graph, while HGL/HGL-QoS consume that native graph directly. Part of the gap above therefore reflects the wider node population and relation set available to the typed model, not relation-specific message passing in isolation. We did not run a native-substrate GL ablation to separate these two effects — doing so would require resolving the collapse the projection substrate exists to avoid — so the $+0.353$ figure should be read as the advantage of the typed pathway as deployed, not as an isolated effect of typing alone.

7.2.1. Empirical Edge-Removal Analysis

We further probed edge criticality by simulating the removal of candidate relationship channels while keeping both endpoint components alive. On the `av_system` topology across 50 candidate bridge edges, 46 edges exhibited zero downstream cascade impact, while the 4 edges with non-zero impact were direct library communication channels (`PUBLISHES_TO` / `SUBSCRIBES_TO`). This demonstrates that individual communication links are largely redundant, whereas component-level failures and shared library dependencies induce disproportionate systemic disruption.

7.3. RQ3: Ablations and Sensitivity Analysis

A note on what ρ measures here. Several sweeps below — AHP shrinkage λ , the topic-weight triple, the QoS sub-weights, the Morris and Dirichlet joint sweeps — report RM’s rank correlation against $I^*(v)$ as their outcome. RM is not proposed as a ranking model (Section 5.1), so within this subsection ρ is used strictly as a *sensitivity probe*: does perturbing a declared constant move a measurable quantity, and by how much relative to the gaps between predictor families elsewhere. A sweep finding one setting higher than another is evidence about that constant’s leverage, not an optimality claim about RM as a ranker.

7.3.1. QoS Feature Encoding

Explicitly encoding continuous QoS attributes in the GNN (HGL-QoS) does not yield a uniform effect — it trades a little in-distribution accuracy for a large out-of-distribution gain. The two directions are summarised below; note that they are measured under different protocols on different fold counts, though both are scored on the Application population.

Table 10: HGL-QoS against HGL under both protocols. Both are Application-scored; the regimes differ in fold count (7 in-distribution scenarios vs. 8 LOSO folds, the latter including the ATM case study).

Protocol	HGL	HGL-QoS	$\Delta\rho$	Folds won	Wilcoxon p
In-distribution (Table 7)	0.725	0.653	-0.072	1/7	0.078 (n.s.)
Inductive LOSO (Table 9)	0.439	0.608	$+0.169$	7/8	0.016

The apparent contradiction between the two tables is a genuine regime effect, not an inconsistency. In-distribution, HGL-QoS *trails* the base typed model, but the deficit is not significant and is small relative to its own scatter (across-scenario $\sigma = 0.35$ against a 0.072 gap), concentrated in two structurally atypical topologies, Hub-and-Spoke (-0.238) and AV (-0.140), with the remaining five within ± 0.084 . This is consistent with QoS features adding redundant signal, and mild overfitting risk, on topologies already seen — the derived DEPENDS_ON graph already embeds much of the same routing and coupling information in its edge weights.

Under LOSO the relationship reverses decisively and *is* significant: $+0.169$, winning 7 of 8 folds. When the topology is unseen, structural degree signatures no longer transfer, whereas a declared QoS contract means the same thing in an unfamiliar architecture as a familiar one — RELIABLE and PERSISTENT carry their semantics across deployments in a way a betweenness percentile does not. QoS encoding is therefore *situational* rather than redundant: insurance against distribution shift, bought at a small, insignificant in-distribution cost. The practitioner rule: HGL when the target resembles the training corpus, HGL-QoS when it does not, and HGL-QoS by default when that is unknown.

7.3.2. Topic-Weight Coefficients (β, α, ψ)

The outer split of Equation 3 is a declared convex combination, not an elicited one, so we measure what rests on it rather than argue for the particular triple. We swept (β, α, ψ) over a grid spanning the whole simplex — including the QoS-only corner $(1, 0, 0)$ and a uniform prior $(\frac{1}{3}, \frac{1}{3}, \frac{1}{3})$ — propagating each point through the derived DEPENDS_ON edge weights into two closed-form scorers, against the revised SizeNorm envelope and re-elicited QoS sub-weights (Section 3.2).

Table 11: Sensitivity of the topic-weight ordering and of downstream rank correlation against $I^*(v)$ to the declared coefficients (β, α, ψ) . Application population, mean over seven scenarios; 375 topics.

(β, α, ψ)	ρ of $w(t)$ vs. shipped	Topo-QoS ρ	RM ρ
$(0.75, 0.15, 0.10)$ <i>shipped</i>	1.000	0.604	0.267
$(0.90, 0.05, 0.05)$	0.966	0.600	0.273
$(0.85, 0.15, 0.00)$	0.950	0.608	0.268
$(1.00, 0.00, 0.00)$	0.852	0.618	0.268
$(0.60, 0.20, 0.20)$	0.970	0.604	0.261
$(0.50, 0.25, 0.25)$	0.942	0.603	0.259
$(\frac{1}{3}, \frac{1}{3}, \frac{1}{3})$ <i>uniform</i>	0.870	0.608	0.256
Spread across grid	—	0.018	0.017

The coefficients are not load-bearing. Across the entire simplex the induced ordering of $w(t)$ never falls below $\rho = 0.852$ against the shipped ordering, and downstream rank correlation moves by at most 0.018 (Topo-QoS) and 0.017 (RM) — an order of magnitude below the gaps between predictor families. Over the 375 corpus topics the QoS term contributes a weighted standard deviation of 0.194, the frequency term 0.021, the payload term 0.020; this corrects an earlier KiB-based SizeNorm divisor of 50 (an unstated ~ 1 EiB envelope) under which the payload term contributed only 0.0039, realizing $\sim 2\%$ of alpha’s declared budget. The claim is that the declared triple drives no reported result, *not* that it is optimal: the QoS-only corner is marginally better for both Topo-QoS (0.618 vs. 0.604) and RM (0.268 vs. 0.267). The inner QoS split $(0.24, 0.62, 0.14)$ is the geometric-mean priority vector of an independently stated Saaty matrix ($CR \approx 0.016$), replacing an earlier matrix solved backward from a target vector; both are pinned by regression tests.

7.3.3. Intra-Dimension AHP Weight Shrinkage

We evaluated the sensitivity of the RM attribution baseline by sweeping a shrinkage parameter $\lambda \in [0, 1]$ blending internal term weights from uniform prior ($\lambda = 0$) to calibrated AHP judgement ($\lambda = 1$). Note that λ shrinks only the *intra-dimension* vectors: the composite weights $(q_R, q_M) = (0.80, 0.20)$ are declared constants and are identical at every λ .

Table 12: Sensitivity of RM rank correlation against $I^*(v)$ to the AHP shrinkage parameter λ , Application population, mean over seven scenarios.

λ Setting	0.00 (Uniform)	0.50	0.70 (Default)	0.80	1.00 (Raw AHP)
Mean Rank Correlation (ρ)	0.348	0.291	0.267	0.256	0.232




Figure 4: Sensitivity of the RM composite’s rank correlation against $I^*(v)$ to the AHP shrinkage parameter λ , Application population, mean over seven scenarios. ρ declines monotonically toward the elicited judgement: the uniform prior at $\lambda = 0$ is the best setting in the sweep.

λ is the most consequential of the ten constants on this probe, and moving it toward the elicited judgement does not raise ρ . ρ declines monotonically from the uniform prior toward elicited judgement (Figure 4): the uniform prior scores highest (0.348 against 0.232 at raw AHP), winning all seven scenarios (paired Wilcoxon, $p = 0.0156$), and the shipped $\lambda = 0.70$ trails it by $\Delta\rho = -0.081$ ($p = 0.0156$), with no plateau around the default. Read as the sensitivity probe this subsection declares — not a ranking-optimality claim about RM (Section 5.1) — the elicited hierarchy buys *transparency*, making the composite auditable and its terms nameable, rather than ranking accuracy. We retain $\lambda = 0.70$ on that basis.

A measurement correction. An earlier version scored $Q(v)$ from features restricted to the Application–Library `DEPENDS_ON` projection rather than the full typed graph. There no Application is an articulation point and no incident edge a bridge in six of eight scenarios, so four of Availability’s five terms vanish and $A(v)$ — carrying $w_R(1 - \alpha) \approx 0.51$ of the composite — is constant, and the sweep returned uniformly negative ρ (-0.146 to -0.111). Figures above use the full pipeline (Section 5.2); both series are retained in the replication artifact. The same correction applies to the threshold and domain-weighting ablations below.

7.3.4. Joint Sensitivity Across All Ten Weight Constants

The sweeps above, like every other ablation here, vary one declared constant at a time. That one-at-a-time (OAT) design cannot detect interactions between factors, or say whether a factor flat in isolation stays flat when a second moves with it [54]. We closed that gap by sweeping all ten hand-set weight constants jointly — (β, α, ψ) , $(w_{\text{rel}}, w_{\text{dur}}, w_{\text{prio}})$, the vertex power-mean exponent p , the library fan-out coefficient γ , the AHP shrinkage λ , and r_α — under two complementary designs on six scenarios (`enterprise_system` excluded on cost grounds, its structural analysis costing roughly $15\times$ any other per evaluation point).

Morris elementary-effects screening [55, 56] takes 10 random trajectories through the full ten-dimensional hyper-rectangle (110 evaluations), deliberately *not* constraining the convex-combination groups to a unit sum, and ranks factors by μ^* , the mean absolute elementary effect on mean ρ against $I^*(v)$: λ and r_α dominate, consistent with the OAT sweeps above. The topic-weight and QoS sub-weight triples sit

Table 13: Morris elementary-effects screening over all ten weight constants, ranked by influence (μ^*) on mean ρ . Six scenarios, 10 trajectories, 110 evaluations.

Factor	μ^*	σ
λ (AHP shrinkage)	0.124	0.077
r_α	0.096	0.035
α	0.019	0.034
w_{rel}	0.019	0.018
β	0.017	0.023
w_{dur}	0.014	0.017
w_{prio}	0.013	0.018
ψ	0.012	0.015
p (power-mean exponent)	0.005	0.005
γ (fan-out)	0.001	0.002

in a comparable, modest band ($\mu^* \in [0.012, 0.019]$) rather than one dominating the other by an order of magnitude — confirming jointly what the corrected term contributions show at the formula level: none of these six factors is individually load-bearing, and none negligible by construction. p and γ are least influential. σ relative to μ^* is largest for λ and α , indicating their effect depends on where the other nine sit — the interaction an OAT sweep cannot diagnose, though within the same order as μ^* rather than dominating it.

Dirichlet simplex sampling asks the sharper question: over the realistic space of declared budgets — respecting the unit-sum constraint on both groups, other factors shipped — how much do accuracy and the shortlist vary. Over 100 draws, mean $\rho = 0.243$ (shipped 0.252), sd = 0.006, range [0.229, 0.260], 90% interval [0.231, 0.253]; rank stability is high (mean Kendall $\tau = 0.899$, worst 0.791; mean top-20% Jaccard 0.825, worst 0.699). The declared budgets are, jointly and on the realistic simplex, not a lever on either the score or the shortlist.

7.3.5. Convergent Validity Across Simulation Oracles

To test construct validity we compared node criticality rankings across three of the four component-level oracles of Section 4.3 — $I^*(v)$ (**FaultInjector**), $I_{\text{comp}}(v)$ (**FailureSimulator**) and $I_{\text{dyn}}(v)$ (**MessageFlowSimulator**) — over seven scenarios and five seeds on the Application population. The Maintainability-only $I_M(v)$ is excluded as it targets a different dimension; its own limitation is discussed in Section 8.3.

Table 14: Inter-oracle agreement. Chance-level top- K Jaccard at $K = 0.2n$ is 0.111; the tie-robust column admits every component tied with the K -th.

Oracle pair	Mean ρ (range)	Mean τ	Jaccard@ K	Tie-robust
I_{dyn} vs. I^*	0.883 (0.756–0.972)	0.745	0.424	0.419
I_{comp} vs. I^*	0.468 (0.171–0.677)	0.349	0.307	0.312
I_{comp} vs. I_{dyn}	0.465 (0.121–0.658)	0.348	0.313	0.313

Ordering converges; the critical set does not. The behavioural queue simulator and the topological cascade oracle agree strongly on ordering ($\rho = 0.883$, never below 0.756). Because the two are built on different principles — delivered message rates under load versus reachability over edges — this is not reducible to a shared construction artifact, and is the strongest convergent evidence that $I^*(v)$ tracks dynamic service disruption rather than only topology. The composite oracle is the outlier ($\rho \approx 0.47$).

Set-level agreement is weaker, and we report it as the binding limitation. Top- K Jaccard is 0.42 for the strongest pair and 0.31 for the others, against 0.111 under independent rankings. Two controls precede any interpretation. It is *not* a tie-breaking artifact: admitting every component tied with the K -th moves it by at most 0.005. Nor is it simply label noise, though noise is a substantial part — the labeler’s own seed-to-seed self-agreement spans 0.44–1.00, so one oracle against *itself* reaches only 0.44 in its worst scenario. Top- K identity at $K = 0.2n$ is intrinsically unstable, and the cross-oracle values sit just below the floor a single

oracle achieves against its own reruns. The defensible claim is about ranking, not membership of a fixed-size set.

A falsified hypothesis, reported. We tested whether weighting both topological oracles by the same $w(t)$ makes them converge — the mechanism by which QoS weighting would be a construct-validity improvement rather than a modelling choice. It does not: $\rho(I_{\text{comp}}, I^*)$ is 0.468 weighted and 0.477 unweighted ($\Delta\rho = -0.009$; Jaccard 0.307 vs. 0.306). We record the negative result rather than leaving the artifact uncited.

Two reasons that null is weaker than it looks. First, the arms are not comparably strong: disabling QoS removes I_{comp} ’s severity term entirely, but on I^* removes only a ladder multiplier that a $[0, 1]$ clamp already suppresses near feed-loss saturation — exactly the high-impact regime a ranking metric is most sensitive to. Toggling that ladder shifts I^* ’s labels by mean absolute 0.064 (range 0.012–0.145) against 0.013 for a shared $w(t)$ -proportional treatment, so the null bounds *this* treatment rather than QoS-aware labelling in general; the untested arm is I_{dyn} ’s delivery-rate label. Second, four of the seven scenarios have $w(t)$ standard deviation ≈ 0.02 (Section 6.1), so a null measured partly on a QoS-flat substrate is expected regardless of treatment strength. We did not separate the two effects, and flag this as an open confound (Section 8.3).

7.3.6. Domain-Specific Weighting Sensitivity

We swept the composite reliability weight w_R over its full range across 10 scenarios, reading off the shipped static default ($w_R = 0.80$), an equal split (0.50), and the domain-derived value. The three are indistinguishable (mean $\rho = 0.353, 0.349, 0.347$), and the sweep explains why: over the entire range mean ρ moves only from 0.341 to 0.365, a total spread of 0.024. The derived w_R lies in $[0.70, 0.76]$ across domains, and mean Kendall fidelity between domain-derived and static rankings is $\tau = 0.974$. This confirms the scoping of Section 5.1: Context-of-Use reweighting is an *attributional* mechanism, not a ranking-improvement device, and should be reported as neither gain nor loss.

7.3.7. Threshold and Normalization Sensitivity

We swept 7 scenarios across a cascade propagation-threshold parameter $\in \{0.0, 0.1, 0.2, 0.35, 0.5, 0.75, 1.0\}$ controlling the oracle’s eligibility cutoff, and separately across 3 label-normalization techniques (robust, min-max, standard) at fixed threshold. Ranking is more sensitive to the threshold ($\Delta\rho = 0.102$; mean ρ rises from 0.177 at 0 to 0.278 at 0.5 then plateaus, the shipped 0.35 sitting inside that plateau at 0.274) than to normalization ($\Delta\rho = 0.022$; robust 0.267 against 0.244 for both alternatives). The threshold is the more consequential free parameter, and the shipped setting sits on the flat part of its curve rather than at a tuned peak. A small spread means the ordering is stable under that parameter — not, by itself, that it is *correct*.

7.3.8. Availability Was Measuring Almost Nothing, and Why

The Availability sub-characteristic $A(v)$ (Section 5) was, in the version these results were first computed on, driven almost entirely by one component of a five-term formula that fires vanishingly rarely here. $AP_c(\text{directed})$, Bridge Ratio and CDI are all articulation-point-gated: a node that is not a Tarjan cut vertex scores exactly 0.0 on all three by construction. On six of eight LOSO scenarios — including the 350-component Enterprise topology — *zero* Application-layer components are articulation points, so all three vanish and $A(v)$ collapses to its remaining $0.05 \cdot w(v)$ QoS term: A_{max} sits at exactly 0.0500 (Enterprise, AV, Microservices, Hub-and-Spoke) or within rounding (ATM 0.0492, Financial 0.0455), with $\sigma_A \in [0.0145, 0.0398]$ against $\sigma_{FT} \in [0.147, 0.171]$. The rule-based SPOF detector, reading the same membership, corroborates this: mean $F_1 = 0.0042$, zero predicted SPOFs in seven of eight scenarios.

This is not evidence that single points of failure are absent, only that a hard cut-vertex test cannot see them in a topology this redundant — on IoT the detector reaches precision 1.0 at recall 0.017, so the one component it flags is genuinely failure-critical and it is blind to the rest. A redundantly but load-bearingly connected component, such as the sole active publisher into a topic with structurally-possible-but-unused alternate routes, fragments nothing when removed yet measurably degrades reachability.

We therefore replaced CDI ’s cut-vertex gate with a continuous redundancy-deficit measure: for every node in the main connected component, the fractional increase in average shortest-path length from a fixed sample

of high-degree sources when v is removed, capped at 1.0 on genuine fragmentation — so a true cut vertex stays at the ceiling (the prior semantics are the limiting case) while a redundant-but-load-bearing node registers a graded score. We rebalanced $A(v)$ ’s AHP weights accordingly, promoting CDI from $a_{CDI} = 0.10$ to parity with the hard term ($a_{CDI} = a_{AP_c} = 0.2563$) since both now measure removal-impact severity at different graduations, and shrinking a_{QSPOF} from 0.25 to 0.1998 to bound its overlap with AP_c . Recomputing moves σ_A to $[0.0254, 0.0530]$, roughly $1.6\text{--}2.0\times$ its prior range, and $A_{\max}$ off the ceiling everywhere (Enterprise: $0.0500 \rightarrow 0.0878$). $A(v)$ is no longer a rescaled QoS weight with zero structural content, though it correctly remains narrower-spread than $FT(v)$: genuine fragility should be the exception.

7.3.9. Anti-Pattern Detection and CI/CD Quality Gates

Validating our rule-based anti-pattern catalog against the composite oracle $I_{\text{comp}}(v)$ yielded mean detection $F_1 = 0.3781$ across 8 scenarios — but F_1 understates what the catalog does: mean precision is 0.239 against mean recall 0.900, so it flags 94.2% of components on average, trading precision for near-exhaustive coverage of true positives. This is a deliberate high-recall CI/CD posture (a missed critical component costs more than a false positive requiring manual triage), and should be read as such rather than via F_1 alone. One detector, DEEP_PIPELINE, is excluded: it enumerates every simple source-to-sink path and does not terminate within ten minutes on the 50-application Healthcare topology — itself a reportable scalability limit of exhaustive path enumeration. The remaining 18 detectors ran in 0.04 s (50 nodes) to 54.85 s (300 nodes), inside standard pre-commit and pull-request budgets.

Stratified against pooled. We stratify the RM composite’s rank correlation by node type: $\rho = 0.503$ (Application), 0.395 (Broker), 0.142 (Node), while the *pooled* correlation across all types is $\rho = 0.028$ — outside the per-type range entirely. This is a Simpson’s-paradox effect: aggregating across populations with different scales and base rates reverses the within-group trend, so the pooled figure is not a summary of the per-type ones, and only the stratified numbers should be quoted when discussing RM’s structural predictive power.

7.3.10. Using RM Correctly: Score Within a Node Type

The stratification above is not only a reporting convention; it is a usage rule, and getting it wrong changes which components an architect is told to fix.

The rule. Score, tier and rank components *within* a node type. Never derive a single ranking or threshold from a population mixing Applications, Brokers, Topics, Infrastructure Nodes and Libraries: their scores differ in scale and base rate, so a shared threshold measures type membership as much as risk. The layer projections of Section 3.3 are the stratification device (π_{infra} , π_{mw} , π_{app}), with only π_{system} spanning all five types; read that one as a structural and visualisation view, and take gating decisions from the single-type layers.

What pooling costs, measured. Classifying every component in the eight-scenario corpus twice — once against one box-plot over the whole system layer, once within its node type — changes the tier of **62.8%** of 1,619 components, and **19.0%** cross the CRITICAL/HIGH boundary that determines whether a component is surfaced at all. The effect is stable across scenarios (tier changes 51–69%, boundary crossings 14–22%), so it is a property of mixing populations rather than of any topology, and its direction is systematic: pooling suppresses the top of whichever type carries the lower score scale. The classifier now derives quartiles and fences within each node type, with types too small for stable quartiles falling back to the pooled fence. The residual case is π_{app} , which still pools Applications with Libraries; their LOSO correlations overlap (0.195 against 0.105), so we retain the pairing as a scope condition.

7.3.11. HGT Attention Weight Analysis

Extracting relational attention matrices from the trained HGT model on the ATM Case Study (Figure 5) revealed that the model places its highest single attention weight ($\alpha_{uv} = 1.00$) on a USES edge (an Application’s dependency on a shared Library), with the next tier ($\alpha_{uv} \approx 0.50$) split between a ROUTES edge (Broker $\rightarrow$ Topic) and further USES/SUBSCRIBES_TO edges. This pattern — dominant attention on library-dependency and broker-routing channels rather than direct publish/subscribe message flow — reflects the shared-library




Figure 5: Relational attention extracted from the trained HGT on the ATM case study. The highest-weighted channels are USES (Application → shared Library) and ROUTES (Broker → Topic) rather than direct publish/subscribe edges, matching the shared-library blast-radius and broker-centrality pathways found by the edge-removal analysis.

blast-radius and broker-centrality failure pathways identified elsewhere in this study (Section 7.2’s edge-removal analysis, Section 3’s simultaneous-failure semantics for shared libraries) rather than sequential pub-sub cascade propagation.

7.4. RQ4: Real-World Distributed Software Architecture Validation

To assess external validity on authentic production architectures, we evaluated SaG on three open-source distributed systems.

Table 15: Empirical validation on authentic real-world distributed software architectures.

Real-World Architecture	$ V $	$ V_{\text{app}} $	Spearman ρ	Kendall τ	$F_1@K$	Tie-Robust F_1	Non-Zero	Gain
Autoware.universe (ROS 2)	75	32	0.688 ± 0.009	0.517	0.800	0.800	19 / 32	+0.36
Cloud Microservices Mesh	60	22	0.778 ± 0.001	0.639	1.000	0.760	8 / 22	+0.01
Train-Ticket Booking Mesh	90	41	0.759 ± 0.001	0.605	1.000	0.810	14 / 41	+0.26

Key Insights for RQ4:

- Strong Real-World Rank Agreement:** SaG achieves high rank correlation on Cloud Microservices ($\rho = 0.778$) and Train-Ticket ($\rho = 0.759$), and solid agreement on Autoware.universe ($\rho = 0.688$).
- Critical-Set Containment:** Every single application with non-zero cascading impact in Cloud Microservices and Train-Ticket is successfully captured within the predicted top- K critical set (tie-robust $F_1@K = 0.760$ and 0.810).
- Substantial Predictive Gain:** SaG outperforms raw degree centrality by +0.360 on Autoware and +0.264 on Train-Ticket, demonstrating that typed dependency derivation captures critical architectural semantics beyond superficial connectivity.

7.5. RQ5: Analysis Cost and CI/CD Deployment Feasibility

RQ5 asks what the analysis costs at CI/CD time and which stage dominates that cost. A learned model invites the question of whether it remains usable as systems grow; on this pipeline the answer is that the neural network is not the constraint — the deterministic structural analysis that feeds it is, by more than two orders of magnitude.

Table 16: Per-stage cost of the deployed (inference) path on generated systems of increasing size, CPU, median of three runs with a warm-up pass excluded. Training is not included: it is a one-off cost paid once over the corpus, not per analysed system.

$ V $	$ E $	Analyse (s)	Graph→tensor (s)	HGT forward (ms)	Analyse : forward
249	1,127	0.27	0.011	22.5	12×
499	2,402	0.95	0.022	15.7	61×
999	6,422	4.74	0.055	36.7	129×
1,998	19,301	23.83	0.153	43.7	545×

The GNN is the cheapest stage. Inference on a 2,000-component system takes 44 ms — a handful of sparse matrix products whose cost grows close to linearly in $|E|$ — while the structural analysis preceding it takes 24 s. The ratio widens with scale rather than narrowing, from 12× at 249 components to 545× at 1,998. Any effort spent making this framework scale should therefore go to the classical graph metrics, not to the model: betweenness is $O(|V||E|)$ and the directed articulation score and closeness are $O(|V|(|V| + |E|))$, giving an overall $O(|V|^2 + |V||E|)$. One mitigation already ships, with the connectivity-degradation term switching to deterministic top-50 core sampling above $|V| = 300$.

Cost tracks edges, not components. The corpus scenarios make this concrete: the 520-component Enterprise mesh takes 27.2 s to analyse while a *denser-in-name-only* 999-component generated system takes 4.7 s, because the former carries 3,245 structural edges against a far sparser fan-out per component. Practitioners sizing a CI/CD budget should estimate from $|E|$, or from $|V| \cdot |E|$, rather than from component count. Within the paper’s corpus the whole gate — analysis plus all 18 evaluated detectors — runs in 0.02 s (29 components) to 27.4 s (520 components), comfortably inside a pull-request budget.

What we have not measured, and do not claim. Nothing above 2,000 components has been timed, and nothing here should be extrapolated past it; all figures are single-threaded CPU, with no GPU measurement; and one detector, DEEP_PIPELINE, is excluded throughout because exhaustive source-to-sink path enumeration does not terminate within ten minutes even at 50 applications (Section 7.3). That detector is the one component of the framework known *not* to scale, and its exclusion is a limitation rather than a tuning choice. Training cost, for completeness: the full seven-variant, eight-fold, five-seed leave-one-scenario-out sweep reported in Table 9 took approximately 45 minutes per learned variant on CPU.

8. Discussion, Threats to Validity, and Conclusion

8.1. Discussion and Practical Implications

Our empirical findings provide clear guidance for software architects and reliability engineers:

- **What the Predictive Pathway Is For:** The typed model earns its place on three counts. It is the best configuration for evaluating an architecture never analysed before ($\rho = 0.608$ out of distribution), where the untyped alternative fails to transfer at all ($\rho = 0.086$). It is the strongest identifier of the critical top- K shortlist hardening actually consumes ($F_1@K = 0.414$). And it answers in 44 ms, letting architectural risk be gated on every pull request rather than swept nightly (Section 7.5). Beyond the ranking it supplies typed attention over the relations that carried a cascade, per-relationship criticality, and multi-task quality outputs — none of which a centrality score produces.
- **Where the Boundary Is — and When Not To Train:** The honest comparison is against a training-free QoS-weighted centrality baseline, and there the ranking margin is +0.037 and not statistically significant ($p = 0.64$, Table 9). A team that needs a critical-component ranking and *nothing else*, on

architectures resembling ones it has already characterised, should use Topo-QoS: it costs no training, no corpus, and no checkpoint. The learned model earns its cost when the additional outputs above are wanted, or when the deployment target differs structurally from anything in the training corpus. Between the two typed variants: HGL for in-distribution accuracy, HGL-QoS otherwise (Section 7.3).

- **What the Explanation Layer Adds:** The deterministic RM profile is what makes a ranked set actionable. Whether a service is critical through single-point-of-failure exposure (Availability) or wide error propagation (Fault Tolerance) dictates whether to add load-balanced replicas or circuit-breaker policies — a distinction neither the predictor nor the oracle expresses.

This split mirrors Figure 1: Pathway A (RM/ $Q(v)$) and Pathway B (Topo-BL/Topo-QoS/GL/HGL) answer different questions from the same graph, so Table 9 and Figure 3 report RM alongside the ranking predictors as a reference point rather than as a competing entry in that comparison.

8.2. Performance and Sustainability Implications

The cost profile of Section 7.5 inverts the usual expectation of a GNN pipeline: the learned model is the *cheapest* stage, and the gap widens with scale (44 ms inference against 23.8 s of structural analysis at 2,000 components; $12\times$ at 249, $545\times$ at 1,998). Two consequences follow. Per-commit gating is affordable — the whole gate runs in 0.02–27.4 s on our corpus, against minutes to hours for the exhaustive simulation it replaces — and effort to scale this framework belongs on the classical graph metrics rather than the network, since betweenness is $O(|V||E|)$ and articulation and closeness $O(|V|(|V| + |E|))$, so budgets follow $|E|$ rather than component count. The sustainability case rests on *avoided* compute: production cascades drive restart loops, failover storms and retransmissions that deliver no user-visible work, and a defect removed at design time never incurs that load. We state its limit plainly, as this is the theme on which this paper is thinnest: **we performed no energy, power, or carbon measurement.** Every figure here is single-threaded CPU wall-clock time, we did not instrument package power or GPU-side inference, and we have not quantified the energy of any avoided failure — which would require observing the production incidents our simulated ground truth stands in for. The claim is therefore conditional, not a measured saving in joules; direct energy attribution is future work (Section 8.4).

8.3. Threats to Validity

- **Construct Validity:** Our ground truth is generated via discrete-event cascade simulation on structural models rather than observing live production outages. To characterise construct divergence we evaluated the three reliability-facing simulation engines (**FaultInjector**, **FailureSimulator**, **MessageFlowSimulator**) of our four-oracle taxonomy (Section 4.3). On *ordering* the evidence is genuinely convergent: the behavioural queue-flow simulator and the topological cascade oracle agree at $\rho = 0.883$, never below 0.756 on any scenario, despite being built on different principles — though that convergence is about *service disruption reach*, not about the operational criticality of the messages lost: the queue-flow simulator’s label is an unweighted delivery-rate delta, blind to message priority and only partly sensitive to reliability policy. On *critical-set membership* it is not: top- K Jaccard is 0.31–0.42 across pairs against 0.111 expected by chance, and this is neither a tie-breaking artifact (≤ 0.005 change under a tie-robust cut) nor purely construct divergence — the labeler compared against its own reruns reaches only 0.44 in its worst scenario, so the statistic is unstable at this K for any oracle. We also could not close the gap by weighting: applying the same $w(t)$ to both topological oracles changes their agreement by $\Delta\rho = -0.009$, a null result we report rather than omit. A further 30–47% of components per scenario carry no simulated ground truth at all, because the cascade model cannot express direct Topic or Node failure. Accordingly, results established against one oracle should not be read as claims about another, and none of them are claims about observed production outages: the reported correlations measure surrogate fidelity to a simulator, not predictive accuracy against deployed systems. The fourth oracle, $I_M(v)$ (**ChangePropagationSimulator**), stands outside this comparison because it targets Maintainability rather than a competing reliability ranking, and it carries a construct-validity limitation of its own worth stating plainly: $M(v)$ (Section 5.2) carries $q_M = 0.20$

of the composite $Q(v)$, yet this paper reports no maintainability correlation, gate, or table against it. The reason is not merely that the labeler never observed it (above) — $I_M(v)$ is computed on every Validate-stage sweep — but that $I_M(v)$ traverses the same derived dependency topology $M(v)$ is scored from, so agreement between them would be an internal consistency check rather than independent evidence. Settling this would require an external referent uncorrelated with that topology, such as code churn or co-change coupling mined from commit history; we did not measure this for our three real-world architectures and leave it to future work.

- **Internal Validity:** Potential feature leakage is prevented by strict graph view separation: predictors operate on G_{analysis} , whereas ground-truth simulators operate on $G_{\text{structural}}$. We additionally identified and fixed a normalization defect in the code-quality scoring pipeline: a min-max helper previously assigned the maximal Code Quality Penalty to any zero-variance population, which is correct for one genuine node but was indistinguishable from “many nodes carrying no code-quality data at all” — the exact shape of every Library in our three real-world scenarios, which carry no source-level `code_metrics`. The fix (scoring an undifferentiated multi-node population as zero rather than maximal penalty) changes Library Maintainability tier classifications in the real-world corpus but, as expected for a fix confined to a population with no variance to rank, leaves Application-population rank correlation against $I^*(v)$ effectively unchanged (Table 15: $\Delta\rho \in \{-0.003, 0.000, 0.000\}$ across the three real-world scenarios). A second, unresolved confound belongs here too: per-topic QoS in four of our seven synthetic scenarios is generated by a domain-keyed lookup whose table collapses to a single entry, so $w(t)$ has standard deviation ≈ 0.02 there against 0.19–0.28 in the other three (Section 6.1). **Topo-QoS** nonetheless gains a mean $+0.305\rho$ over **Topo-BL** on exactly those four scenarios (AV, Enterprise, Hub-and-Spoke, Microservices; Table 7), a gain that cannot be QoS discrimination given the absent signal, and must instead originate in the weighting scheme’s other effects — for instance its uniform rescaling of pub/sub-derived `DEPENDS_ON` edge weights relative to `USES`-derived ones. We did not run the diagnostic that would separate the two (re-scoring **Topo-QoS** with $w(t)$ forced constant), so we report the confound rather than an explanation for it; the same corpus property is the more likely driver of Section 7.3’s $\Delta\rho = -0.009$ QoS-convergence null than the ladder clipping we attribute it to there.
- **External Validity:** While our corpus spans ten architectures across six declared system domains, three of them authentic open-source systems (ROS 2 and microservices), future work should evaluate larger enterprise deployments with hundreds of microservices.
- **Sustainability Claims Are Unmeasured:** The efficiency argument in Section 8.2 rests on wall-clock cost and on avoided recovery compute, neither of which we instrumented for energy. No power, energy or carbon figure appears in this paper, and none should be inferred from the timing tables; a joules-per-analysis claim, or a claim about energy saved per prevented outage, would require hardware counters and a live incident testbed we did not have.
- **Conclusion Validity:** Given the heavy-tailed, non-normal distribution of cascading failure impacts, all statistical comparisons utilize non-parametric rank correlation (Spearman ρ , Kendall τ), bootstrap confidence intervals ($B = 2,000$), and paired Wilcoxon signed-rank tests. Two hazards deserve explicit statement because both changed conclusions in this study rather than merely threatening to.

First, *aggregation across node types*. The RM composite’s rank correlation pooled across types ($\rho = 0.028$) falls outside the range spanned by its own per-type correlations ($\rho \in [0.14, 0.50]$) — a Simpson’s-paradox effect. Reported pooled, the same LOSO experiment placed the RM baseline at $\rho = -0.014$ and the untyped GL model at 0.381; reported on the Application population it places them at $+0.195$ and 0.086 respectively, reversing their order. Every figure in this paper is therefore scored on a single stated population (Section 6.3), and we quote no pooled cross-type correlation.

Second, *substrate*. An earlier version of the RM ablations in Section 7.3 scored $Q(v)$ from features restricted to the Application–Library `DEPENDS_ON` projection — the substrate built for GNN feature/label alignment — on which no Application is an articulation point and no incident edge is a bridge in six

of eight scenarios. Four of Availability’s five terms vanish there, leaving $A(v)$ (roughly 0.51 of the composite) constant, and all three affected sweeps consequently reported negative ρ for a baseline that is in fact weakly positive. The ablations are now computed through the full analysis pipeline. We record this because the failure mode is not visible in the output: a degenerate score produces plausible-looking correlations, and only checking the variance of each term exposed it.

8.4. Limitations and Future Work

1. **Safety and Security Integration:** SaG currently focuses on Reliability and Maintainability. Incorporating formal hazard categories (e.g., ISO 26262 ASIL ratings) and security threat models into the multigraph schema represents an important direction for future research.
2. **Hardware-in-the-Loop (HIL) Validation:** Validating SaG’s predictions against real-time physical fault-injection testbeds in cyber-physical environments.
3. **Direct Energy Attribution:** Instrumenting the pipeline with hardware power counters, measuring GPU-side inference, and pairing predicted blast radii against measured recovery cost on a live testbed, so that the avoided-compute argument of Section 8.2 can be stated in energy rather than in wall-clock time.
4. **Automated Architectural Refactoring:** Extending SaG from predictive analysis to prescriptive synthesis—automatically generating pull requests that reconfigure QoS policies and add redundancy.

8.5. Conclusion

We presented **Software-as-a-Graph (SaG)**, a pre-deployment Static System Analysis framework whose primary instrument is a relation-specific Heterogeneous Graph Transformer that forecasts cascading failure blast radii from Architecture-as-Code manifests alone, before any telemetry exists, and whose predictions are made actionable by an interpretable ISO-grounded Reliability–Maintainability explanation layer. Together they bridge the Architecture–Code Gap.

Our empirical results across synthetic and authentic open-source distributed systems show that relation-specific typing is what makes graph learning transfer to unseen architectures: the typed model reaches $\rho = 0.608$ where its untyped counterpart collapses to 0.086, an advantage significant across all eight folds, and it identifies the critical top- K set better than any alternative we evaluated ($F_1@K = 0.414$) at a cost of 44 ms per system — two orders of magnitude below the structural analysis feeding it, which is what makes per-commit architectural gating practical. We are equally explicit about the boundary of that result: against a training-free QoS-weighted structural baseline the typed model’s ranking advantage is +0.037 and not statistically significant, so the case for graph learning here rests on transfer to unfamiliar architectures and on what it supplies beyond a ranking — typed attention, multi-task quality outputs, relationship-level criticality — rather than on ranking accuracy alone.

CRediT authorship contribution statement

[Omitted for double-anonymised review. To be completed on acceptance with the standard CRediT roles: Conceptualization; Methodology; Software; Validation; Formal analysis; Investigation; Data curation; Writing – original draft; Writing – review and editing; Visualization; Supervision.]

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Funding

[Omitted for double-anonymised review.]

Data availability

The replication package—including the seven synthetic scenario datasets and the configurations that generate them, the topology generator, the four simulation harnesses, the real-world architecture adapters, and every analysis script behind the reported tables and figures—will be made openly available upon publication. The synthetic corpus is regenerable rather than merely archived: each dataset carries its seed and a SHA-256 digest in a committed manifest, and a regression test asserts byte-identical regeneration from the configurations (Section 6.1.1). Every table and figure in this paper is produced from a committed artifact by a script in that package; none of the reported values is transcribed by hand.

Declaration of generative AI use

[To be completed by the authors in accordance with the journal’s policy.]

References

- [1] S. Macenski, T. Foote, B. Gerkey, C. Lalancette, W. Woodall, Robot operating system 2: Design, architecture, and uses in the wild, *Science Robotics* 7 (66) (2022) eabm6074.
- [2] J. Kreps, N. Narkhede, J. Rao, Kafka: A distributed messaging system for log processing, in: *Proc. 6th Int. Workshop on Networking Meets Databases (NetDB)*, 2011.
- [3] Object Management Group, Data distribution service (dds), Tech. Rep. formal/2015-04-10, version 1.4, Object Management Group (2015).
- [4] OASIS, MQTT version 5.0, OASIS Standard (2019).
- [5] P. T. Eugster, P. A. Felber, R. Guerraoui, A.-M. Kermarrec, The many faces of publish/subscribe, *ACM Computing Surveys* 35 (2) (2003) 114–131.
- [6] International Organization for Standardization, ISO/IEC 25010:2023 — systems and software engineering — systems and software quality requirements and evaluation (square) — product quality model, Tech. rep., International Organization for Standardization (2023).
- [7] International Organization for Standardization, ISO/IEC 25019:2023 — systems and software engineering — systems and software quality requirements and evaluation (square) — quality-in-use model, Tech. rep., International Organization for Standardization (2023).
- [8] T. J. McCabe, A complexity measure, *IEEE Transactions on Software Engineering* SE-2 (4) (1976) 308–320.
- [9] S. R. Chidamber, C. F. Kemerer, A metrics suite for object oriented design, *IEEE Transactions on Software Engineering* 20 (6) (1994) 476–493.
- [10] N. Fenton, J. Bieman, *Software Metrics: A Rigorous and Practical Approach*, 3rd Edition, CRC Press, 2014.
- [11] A. Basiri, N. Behnam, R. de Rooij, L. Hochstein, L. Kosewski, J. Reynolds, C. Rosenthal, Chaos engineering, *IEEE Software* 33 (3) (2016) 35–41.
- [12] L. C. Freeman, A set of measures of centrality based on betweenness, *Sociometry* 40 (1) (1977) 35–41.
- [13] S. Brin, L. Page, The anatomy of a large-scale hypertextual web search engine, *Computer Networks and ISDN Systems* 30 (1–7) (1998) 107–117.
- [14] U. Brandes, A faster algorithm for betweenness centrality, *Journal of Mathematical Sociology* 25 (2) (2001) 163–177.

- [15] M. E. J. Newman, *Networks: An Introduction*, Oxford University Press, 2010.
- [16] T. N. Kipf, M. Welling, Semi-supervised classification with graph convolutional networks, in: *Proc. Int. Conf. on Learning Representations (ICLR)*, 2017.
- [17] W. L. Hamilton, R. Ying, J. Leskovec, Inductive representation learning on large graphs, in: *Advances in Neural Information Processing Systems 30 (NeurIPS)*, 2017, pp. 1024–1034.
- [18] P. Veličković, G. Cucurull, A. Casanova, A. Romero, P. Liò, Y. Bengio, Graph attention networks, in: *Proc. Int. Conf. on Learning Representations (ICLR)*, 2018.
- [19] [Anon-A], Authors’ prior work on multi-layer graph dependency analysis for publish–subscribe systems, citation withheld for double-anonymised review (n.d.).
- [20] V. R. Basili, L. C. Briand, W. L. Melo, A validation of object-oriented design metrics as quality indicators, *IEEE Transactions on Software Engineering* 22 (10) (1996) 751–761.
- [21] N. Nagappan, T. Ball, Static analysis tools as early indicators of pre-release defect density, in: *Proc. 27th Int. Conf. on Software Engineering (ICSE)*, 2005, pp. 580–586.
- [22] T. Zimmermann, R. Premraj, A. Zeller, Predicting defects for Eclipse, in: *Proc. 3rd Int. Workshop on Predictor Models in Software Engineering (PROMISE)*, 2007.
- [23] T. Menzies, J. Greenwald, A. Frank, Data mining static code attributes to learn defect predictors, *IEEE Transactions on Software Engineering* 33 (1) (2007) 2–13.
- [24] J. Garcia, D. Popescu, G. Edwards, N. Medvidovic, Toward a catalogue of architectural bad smells, in: *Proc. 5th Int. Conf. on the Quality of Software Architectures (QoSA)*, LNCS 5581, 2009, pp. 146–162.
- [25] J. Garcia, D. Popescu, G. Edwards, N. Medvidovic, Identifying architectural bad smells, in: *Proc. 13th European Conf. on Software Maintenance and Reengineering (CSMR)*, 2009, pp. 255–258.
- [26] D. Taibi, V. Lenarduzzi, On the definition of microservice bad smells, *IEEE Software* 35 (3) (2018) 56–62.
- [27] N. Dragoni, S. Giallorenzo, A. L. Lafuente, M. Mazzara, F. Montesi, R. Mustafin, L. Safina, *Microservices: Yesterday, today, and tomorrow*, in: *Present and Ulterior Software Engineering*, Springer, 2017, pp. 195–216.
- [28] W. Cunningham, The WyCash portfolio management system, in: *Addendum to the Proc. Conf. on Object-Oriented Programming Systems, Languages, and Applications (OOPSLA)*, 1992, pp. 29–30.
- [29] Z. Li, P. Avgeriou, P. Liang, A systematic mapping study on technical debt and its management, *Journal of Systems and Software* 101 (2015) 193–220.
- [30] J. Humble, D. Farley, *Continuous Delivery: Reliable Software Releases through Build, Test, and Deployment Automation*, Addison-Wesley, 2010.
- [31] L. Chen, Continuous delivery: Huge benefits, but challenges too, *IEEE Software* 32 (2) (2015) 50–54.
- [32] International Organization for Standardization, ISO/IEC 25023:2016 — systems and software engineering — systems and software quality requirements and evaluation (square) — measurement of system and software product quality, Tech. rep., International Organization for Standardization (2016).
- [33] International Organization for Standardization, ISO/IEC 25021:2012 — systems and software engineering — systems and software quality requirements and evaluation (square) — quality measure elements, Tech. rep., International Organization for Standardization (2012).

- [34] T. L. Saaty, The Analytic Hierarchy Process: Planning, Priority Setting, Resource Allocation, McGraw-Hill, 1980.
- [35] R. Albert, H. Jeong, A.-L. Barabási, Error and attack tolerance of complex networks, *Nature* 406 (2000) 378–382.
- [36] A. E. Motter, Y.-C. Lai, Cascade-based attacks on complex networks, *Physical Review E* 66 (2002) 065102(R).
- [37] S. V. Buldyrev, R. Parshani, G. Paul, H. E. Stanley, S. Havlin, Catastrophic cascade of failures in interdependent networks, *Nature* 464 (2010) 1025–1028.
- [38] C. Fan, L. Zeng, Y. Sun, Y.-Y. Liu, Finding key players in complex networks through deep reinforcement learning, *Nature Machine Intelligence* 2 (2020) 317–324.
- [39] C. Fan, L. Zeng, Y. Ding, M. Chen, Y. Sun, Z. Liu, Learning to identify high betweenness centrality nodes from scratch: A novel graph neural network approach, in: *Proc. 28th ACM Int. Conf. on Information and Knowledge Management (CIKM)*, 2019, pp. 559–568.
- [40] A. Varbella, K. Amara, M. El-Assady, B. Gjorgiev, G. Sansavini, PowerGraph: A power grid benchmark dataset for graph neural networks, in: *Advances in Neural Information Processing Systems 37 (NeurIPS 2024)*, Datasets and Benchmarks Track, 2024, arXiv:2402.02827.
- [41] M. Schlichtkrull, T. N. Kipf, P. Bloem, R. van den Berg, I. Titov, M. Welling, Modeling relational data with graph convolutional networks, in: *Proc. European Semantic Web Conference (ESWC)*, 2018, pp. 593–607.
- [42] X. Wang, H. Ji, C. Shi, B. Wang, Y. Ye, P. Cui, P. S. Yu, Heterogeneous graph attention network, in: *Proc. The Web Conference (WWW)*, 2019, pp. 2022–2032.
- [43] Z. Hu, Y. Dong, K. Wang, Y. Sun, Heterogeneous graph transformer, in: *Proc. The Web Conference (WWW)*, 2020, pp. 2704–2710.
- [44] X. Fu, J. Zhang, Z. Meng, I. King, MAGNN: Metapath aggregated graph neural network for heterogeneous graph embedding, in: *Proc. The Web Conference (WWW)*, 2020, pp. 2331–2341.
- [45] Z. Ying, D. Bourgeois, J. You, M. Zitnik, J. Leskovec, GNNExplainer: Generating explanations for graph neural networks, in: *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 32, 2019, pp. 9244–9255.
- [46] F. Xia, T.-Y. Liu, J. Wang, W.-S. Zhang, H. Li, Listwise approach to learning to rank: Theory and algorithm, in: *Proc. 25th Int. Conf. on Machine Learning (ICML)*, 2008, pp. 1192–1199.
- [47] Team SimPy, Simpy: Event discrete simulation for Python, <https://simpy.readthedocs.io> (2020).
- [48] S. Kato, S. Tokunaga, Y. Maruyama, S. Maeda, M. Hirabayashi, Y. Kitsukawa, A. Monrroy, T. Ando, Y. Fujii, T. Azumi, Autoware on board: Enabling autonomous vehicles with embedded systems, in: *Proc. ACM/IEEE 9th Int. Conf. on Cyber-Physical Systems (ICCPS)*, 2018, pp. 287–296.
- [49] Google Cloud Platform, Online boutique: A cloud-native microservices demo application, *Software artifact* (2024).
- [50] X. Zhou, X. Peng, T. Xie, J. Sun, C. Ji, W. Li, D. Ding, Fault analysis and debugging of microservice systems: Industrial survey, benchmark system, and empirical study, *IEEE Transactions on Software Engineering* 47 (2) (2021) 243–260.
- [51] F. Wilcoxon, Individual comparisons by ranking methods, *Biometrics Bulletin* 1 (6) (1945) 80–83.

- 1256 [52] B. Efron, R. J. Tibshirani, An Introduction to the Bootstrap, Chapman & Hall, 1993.
- 1257 [53] C. Spearman, The proof and measurement of association between two things, American Journal of
1258 Psychology 15 (1) (1904) 72–101.
- 1259 [54] A. Saltelli, M. Ratto, T. Andres, F. Campolongo, J. Cariboni, D. Gatelli, M. Saisana, S. Tarantola,
1260 Global Sensitivity Analysis: The Primer, John Wiley & Sons, 2008.
- 1261 [55] M. D. Morris, Factorial sampling plans for preliminary computational experiments, Technometrics 33 (2)
1262 (1991) 161–174.
- 1263 [56] F. Campolongo, J. Cariboni, A. Saltelli, An effective screening design for sensitivity analysis of large
1264 models, Environmental Modelling & Software 22 (10) (2007) 1509–1518.